

Python Data Science

Personal Notes

Patrick Bucher

2019-02-10

Contents

1	IPython	2
1.1	Help	2
1.2	Readline Commands	3
1.2.1	Navigation	3
1.2.2	Manipulation	3
1.2.3	History	3
1.2.4	Miscellaneous	3
1.3	Magic Commands	3
1.4	History	4
1.5	Shell Interaction	4
1.6	Miscellaneous	4
1.7	Debugging	5
1.8	Timing and Profiling	5
1.8.1	Timing	5
1.8.2	Runtime Profiling	5
1.8.3	Memory Profiling	5
2	NumPy	6
2.1	Array Creation	6
2.1.1	Arrays of Python Lists	6
2.1.2	Arrays from Scratch	6
2.1.3	Data Types	7
2.2	Array Manipulation	7
2.2.1	Attributes	7
2.2.2	Indexing	8
2.2.3	Slicing	9
2.2.4	Reshaping	10
2.2.5	Concatenation	11
2.2.6	Splitting	12
2.3	Universal Functions	13
2.3.1	Common UFuncs	13
2.3.2	Advanced Features	14
2.4	Aggregations	15
2.5	Broadcasting	16
2.6	Boolean Arrays	18
2.6.1	Bitmasks	19
2.7	Fancy Indexing	20
2.7.1	Broadcasting	20

2.7.2 Assignment	21
2.8 Sorting	22
2.9 Structured Arrays	23
2.10 Date and Time	24
3 Pandas	25
3.1 Series	25
3.1.1 Creation	26
3.1.2 Access: Indexing and Selection	27
3.1.3 Explicit and Implicit Indexing	28
3.2 DataFrame	28
3.2.1 Creation	29
3.2.2 Access: Indexing and Selection	30
3.3 Index	32
3.4 Operations	33
3.5 Handling Missing Data	34
3.6 Hierarchical Indexing	37
3.6.1 Creation of Hierarchical Indices	38
3.6.2 Indexing and Slicing	40
3.6.3 Rearranging Multi-Indices	42
3.6.4 Multi-Indices vs. Panels	44
3.7 Combining Datasets	44
3.7.1 Concat and Append	45
3.7.2 Merge and Join	47
3.8 Aggregation	53
3.9 Grouping	54
3.10 Pivot Tables	57
3.11 Vectorized String Operations	59
3.12 Time Series	61
3.12.1 Sequences	63
3.12.2 Resampling, Shifting, Windowing	65
3.13 High-Performance Pandas: <code>eval()</code> and <code>query()</code>	71
3.14 Miscellaneous	72
4 Matplotlib	73
4.1 Interfaces: MATLAB-style and Object Oriented	76
4.2 Line Plots	79
4.3 Limits, Labels, Legends	81
4.4 Scatter Plots	82
4.5 Visualizing Errors	87
4.6 Density and Contour Plots	91
4.7 Histograms, Binnings and Density	99

1 IPython

1.1 Help

- `help([symbol])` or `[symbol]?:` display the docstring of the symbol
 - Example: `help(map)` or `map?`
- `[symbol]??:` display the source code of the symbol (only if written in Python)
- `<Tab>`-completion: display matching `dir()` entries
- `*` (wildcard): matches any (also empty) string

1.2 Readline Commands

- C means Ctrl
- M means Alt

1.2.1 Navigation

- C-a: move to beginning of the line
- C-e: move to end of the line
- C-f: move one character forward
- C-b: move one character backward
- A-f: move one word forward
- A-b: move one word backward

1.2.2 Manipulation

- C-d: delete character under the cursor
- A-d: delete rest of the word under the cursor (right side)
- C-k: delete to the end of the line (right side)
- C-u: delete the beginning of the line (left side)
- C-y: yank (paste) text deleted before
- C-t: transpose; move character under the cursor one position to the left

1.2.3 History

- C-p: previous command (type multiple times to move back through the history)
- C-n: next command (type multiple times to move forth through the history)
- C-r: search backward in history

1.2.4 Miscellaneous

- C-l: clear screen
- C-c: cancel current command
- C-d: terminate session

1.3 Magic Commands

- %paste: paste code from the clipboard
- %cpaste: paste multiple code snippets interactively, end with --
- %run: run a script and keep the loaded symbols in the REPL
- %history: display the command history
 - %history -n 1-4: display from the first to the fourth command
- %rerun: run a part of the history again
- %save: store the history in a file
- %lsmagic: list magic functions
- %xmode: set exception reporting mode
 - Plain: most compact, least information
 - Context: more information

- Verbose: most detailed output
- %load_ext: load the extension with the given name
- %%file/%%writefile: write the following code section to a file with the given file name
 - a for appending instead of overwriting

To get help on a magic command, use the question mark notation as with any other command. Example: %rerun? shows the documentation for the %rerun magic command.

- %automagic: toggle automagic setting

If %automagic is set, shell commands like cat, cp, env, ls, man, mkdir, more, mv, pwd, rm, rmdir can be used without prefixes. Otherwise, a % prefix is needed.

1.4 History

Lines of input and output are numbered so that single lines can be addressed:

- In: list of all inputs
 - In[4]: fourth input line
- Out: map of all outputs
 - Out[2]: second output line
- _ (single underscore): last output
- __ (double underscore): second to the last output
- ___ (triple underscore): third to the last output
- _n (single underscore with number): n to the last output `_4 = Out[4]`

1.5 Shell Interaction

- ! at the beginning of a line: execute a shell command
- files = !ls -l: store output of a shell command as a list
 - files.grep('foo'): filter list by 'foo'
 - files.fields(1, 2): display columns 1 and 2 of the output
- !mkdir {folder}: create a directory with the variable folder's value as a name
 - surround a Python variable with curly braces to make it available for the shell

1.6 Miscellaneous

- ; at the end of a line: suppress output

1.7 Debugging

Python's standard debugger is `pdb`. IPython comes with an enhanced version `ipdb`.

- `%debug`: start a debugging session starting from the last exception
- `%pdb on`: start debugging session automatically when an exception occurs

Debugging sessions have special commands (usually, only the first letters needs to be typed):

- `l(list)`: show the current location in the file
- `u(p)/d(own)`: move up and down in the call stack
- `n(ext)`: execute current line and move to next line (step over)
- `s(tep)`: enter the function (step in)
- `r(eturn)`: leave the function (step out)
- `q(uit)`: leave the debugging session and exit the program execution
- `c(ontinue)`: leave the debugging session, but keep the program running
- `<Enter>`: repeat previous command
- `p(rint)`: print variables
- `h(help)`: display a list of all available commands or help to the command argument supplied

1.8 Timing and Profiling

1.8.1 Timing

- `%time`: measure the execution time of a single statement/function call
 - The garbage collector will be deactivated so that the result is not biased.
- `%timeit`: measure the average execution time of a single statement/function call after repeated runs
 - The number of runs will be determined automatically.
- `%%timeit`: as above, but working on whole sections of code

1.8.2 Runtime Profiling

- `%prun`: runtime profile of a single statement/function call using Python's built-in profiler
- `%lprun`: line by line runtime profile of a single statement/function call
 - install with `pip install line_profiler` on the shell
 - load with `%load_ext line_profiler` in IPython

1.8.3 Memory Profiling

- install with `pip install memory_profiler` on the shell
- load with `%load_ext memory_profiler` in IPython
- `%memit`: memory profile of a single statement/function call
- `%mprun`: line by line memory profile of a single function call

`%mprun` requires the profiled code to be in it's own module. Example session:

```
%load_ext memory_profiler
%%file fibonacci.py
def fib(n):
    if n == 1 or n == 2:
        return 1
    return fib(n-1) + fib(n-2)

from fibonacci import fib
%mprun -f fib fib(35)
```

2 NumPy

Arrays of numbers are the fundamental data structure for data analysis. Python's primitive values have a large overhead. This information is redundant in lists, because the same type information is stored for every element. NumPy arrays are much more efficient than Python's lists—especially for big data sets. Python also offers an array type without redundant type information. However, this array type doesn't offer the fast and powerful operations of NumPy's ndarray type.

Conventionally, the NumPy library is imported as follows:

```
import numpy as np
```

2.1 Array Creation

2.1.1 Arrays of Python Lists

NumPy arrays can be created from Python lists:

```
>>> ints = np.array([2, 4, 6, 8]) # integer array
>>> floats = np.array([2, 4, 6, 8.1]) # upcast to float because of 8.1
>>> floats = np.array([2, 4, 6, 8], dtype='float') # with explicit type parameter
>>> ints = np.array([1.1, 2.2, 3.3], dtype='int') # with explicit type parameter
```

NumPy arrays can be multi-dimensional:

```
matrix = np.array([[1, 2, 3],
                   [4, 5, 6],
                   [7, 8, 9]])
```

2.1.2 Arrays from Scratch

Numpy offers various functions to generate arrays from scratch. Where a dimension is required (size), a single number (length), a tuple of two (rows, columns) or more (1st dimension, 2nd dimension, 3rd dimension, etc.) can be passed.

- `np.zeros(size, dtype)`: array of zeros
- `np.ones(size, dtype)`: array of ones
- `np.full(size, value)`: array filled with the given value
- `np.arange(start, end, step)`: array with values from start (inclusive) to end (exclusive) and given step width; `length=(end-start)/step`
- `np.linspace(from, to, n)`: array with evenly spaced values in interval `[from,to]` (both inclusive) of length `n`

- `np.random.random(size)`: uniformly distributed random values
- `np.random.normal(mean, sd, size)`: normally distributed array with the given mean and standard deviation
- `np.random.randint(from, to, size)`: random integers in the interval `[from, to)` (inclusive/exclusive)
- `np.random.choice(a, size, replace, p)`: random values from the array `a` or up to the upper bound value `a` with (`replace=True`) or without (`replace=False`) replacement and an optional array of probabilities `p`
- `np.eye(n)`: identity matrix with `n` rows and columns (values at indices with equal row/column index are 1)
- `np.empty(size)`: uninitialized array, values from current memory content (garbage)

2.1.3 Data Types

The `dtype` parameter can either be passed as a string literal or using a pre-defined constant:

1. literal: `dtype='int32'`
2. constant: `dtype=np.int32`

Common numeric types are:

- boolean: `Bool_`
- signed integers: `int8`, `int16`, `int32`, `int64`
 - `int_`: system's default long
 - `intc`: system's default int
- unsigned integers: `uint8`, `uint16`, `uint32`, `uint64`
- floating point: `float16`, `float32`, `float64`
 - `float_`: system default
- complex numbers: `complex64`, `complex128`
 - `complex_`: system default

2.2 Array Manipulation

NumPy arrays offer a rich set of attributes and operation for their manipulation. Since NumPy arrays are the foundation of many higher-level libraries, data manipulation in Python is often NumPy array manipulation.

2.2.1 Attributes

These read-only attributes can be used to retrieve information about an array:

- `ndim`: number of dimensions
- `shape`: size of each dimension
- `size`: total size of the array (the number of elements)
- `dtype`: data type of the array's elements
- `itemsize`: byte size of a single element
- `nbytes`: byte size of the entire array

In general, `nbytes` is equal to `itemsize` multiplied by `size`.

```

>>> np.random.seed(0) # for reproducible results
>>> arr = np.random.randint(10, 100, (3, 3))
>>> arr
array([[54, 57, 74],
       [77, 77, 19],
       [93, 31, 46]])
>>> arr.ndim
2
>>> arr.shape
(3, 3)
>>> arr.size
9
>>> arr.dtype
dtype('int64')
>>> arr.itemsize
8
>>> arr.nbytes
72
>>> arr.itemsize * arr.size
72

```

2.2.2 Indexing

Values of NumPy arrays can both be retrieved and modified by the means of indexing.

The indexing of single dimension arrays works with square brackets, just like indexing of Python lists:

- `arr[0]`: first element
- `arr[n]`: nth element
- `arr[-1]`: last element (first element counted from the end)
- `arr[-3]`: third last element (third element counted from the end)

For multi dimension arrays, a comma separated tuple has to be passed in square brackets:

- `arr[0, 0]`: first element of the first dimension
- `arr[3, 5]`: fifth element of the third dimension

```

>>> np.random.seed(0) # for reproducible results
>>> arr = np.random.randint(10, 100, (3, 3))
>>> arr
array([[54, 57, 74],
       [77, 77, 19],
       [93, 31, 46]])

>>> arr[0, 0]
54

>>> arr[1, 2]
19

>>> arr[-1, -1]
46

```



```
>>> arr[2, 2]
46
```

2.2.3 Slicing

The slicing syntax of Python lists also works for NumPy arrays:

- [start:stop:step], with values omitted defaulting to:
 - start=0
 - stop=[size of dimension]
 - step=1
- For a negative step size, the defaults for start and stop are swapped.

```
>>> arr = np.arange(1, 10)
>>> arr
array([1, 2, 3, 4, 5, 6, 7, 8, 9])

>>> arr[2:5] # third (inclusive) to fifth (exclusive)
array([3, 4, 5])

>>> arr[::2] # every other (beginning with first)
array([1, 3, 5, 7])

>>> arr[::-1] # reversed
array([9, 8, 7, 6, 5, 4, 3, 2, 1])
```

If a step is indicated, two colons are required. Otherwise, step is interpreted as the stop.

Multi-dimension arrays can be sliced by providing multiple, comma-separated slices:

- [start1:stop1:step1, start2:stop2:step2], for slicing the first and second dimension.
- Indexing and slicing can be combined in order to access individual columns/rows:
 - [:, 0]: all rows, first column
 - [0, :]: first row, all columns
 - * [0]: shorthand (: can be omitted)

```
>>> np.random.seed(0) # for reproducible results
>>> arr = np.random.randint(10, 100, (3, 3))
>>> arr
array([[54, 57, 74],
       [77, 77, 19],
       [93, 31, 46]])

>>> arr[::2, 0:2] # columns 0 and 1 of every other row
array([[54, 57],
       [93, 31]])

>>> arr[:, 0] # first column
array([54, 77, 93])

>>> arr[0, :] # first row
array([54, 57, 74])
```

```
>>> arr[0] # first row (shorthand)
array([54, 57, 74])
```

Unlike Python lists, slices of NumPy arrays are *views* to the original data, not *copies* of it. To get a copy of a slice that can be modified without affecting the underlying array, the `copy()` method can be used. Using the array from above:

```
>>> s = arr[:,2, 0:2] # view on columns 0 and 1 of every other row
>>> s
array([[54, 57],
       [93, 31]])
```

```
>>> s[0,1] = 88
>>> s[1,0] = 99
>>> s
array([[54, 88],
       [99, 31]])
```

```
>>> t = arr[1, 0:2].copy() # copy of columns 0 and 1 of the second row
>>> t
array([77, 77])
```

```
>>> t[0] = 11
>>> t[1] = 22
>>> t
array([11, 22])
```

```
>>> arr
array([[54, 88, 74], # 88 introduced through s
       [77, 77, 19], # 11 and 22 missing (working on copy t)
       [99, 31, 46]]) # 99 introduced through s
```

2.2.4 Reshaping

There are two options to reshape an existing array:

1. The function `reshape(size)`, which reshapes the underlying array to the given size (dimension indications).
 - The new size must match the array's size.
 - Good: `arr.size=60, arr.reshape((6, 10))`, because $6*10=60$
 - Bad: `arr.size=16, arr.reshape((4, 6))`, because $4*6>16$
2. Using the slicing parameter `np.newaxis`, which converts a one-dimensional to a two-dimensional array.
 - `arr[np.newaxis, :]`: array elements as columns
 - `arr[:, np.newaxis]`: array elements as rows

```
>>> np.arange(1, 10).reshape((3, 3))
array([[1, 2, 3],
       [4, 5, 6],
       [7, 8, 9]])
```

```
>>> np.arange(1, 4)[np.newaxis, :]
array([[1, 2, 3]])
```

```
>>> np.arange(1, 4)[: , np.newaxis]
array([[1],
       [2],
       [3]])
```

2.2.5 Concatenation

The options to concatenate arrays of same and different dimensions are:

1. The function `np.concatenate(arrays, axis)`, which works on arrays of the same dimensions.
 - `arrays`: a list or tuple of arrays
 - `axis`: index of the axis, along which the concatenation takes place (0: rows, 1: columns, 2: third dimension)
2. Functions, which concatenate the given arrays of (possible) different dimensions:
 - `np.vstack(arrays)`: stack the arrays vertically
 - `np.hstack(arrays)`: stack the arrays horizontally
 - `np.dstack(arrays)`: stack the arrays along the third dimension

```
>>> a = np.arange(1, 5) # 1, 2, 3, 4
>>> b = np.arange(5, 9) # 5, 6, 7, 8
>>> np.concatenate((a, b))
array([1, 2, 3, 4, 5, 6, 7, 8])
```

```
>>> x = a.reshape((2, 2))
>>> x
array([[1, 2],
       [3, 4]])
```

```
>>> y = b.reshape((2, 2))
>>> y
array([[5, 6],
       [7, 8]])
```

```
>>> np.concatenate((x, y), axis=0) # along rows
array([[1, 2],
       [3, 4],
       [5, 6],
       [7, 8]])
```

```
>>> np.vstack((x, y)) # same, but shorter
array([[1, 2],
       [3, 4],
       [5, 6],
       [7, 8]])
```

```
>>> np.concatenate((x, y), axis=1) # along columns
array([[1, 2, 5, 6],
       [3, 4, 7, 8]])
```

```
>>> np.hstack((x, y)) # same, but shorter
array([[1, 2, 5, 6],
       [3, 4, 7, 8]])
```

```

>>> i = np.arange(1, 4).reshape((3, 1))
>>> i
array([[1],
       [2],
       [3]])

>>> j = np.arange(4, 10).reshape(3, 2)
>>> j
array([[4, 5],
       [6, 7],
       [8, 9]])

>>> np.hstack((i, j))
array([[1, 4, 5],
       [2, 6, 7],
       [3, 8, 9]])

>>> m = np.arange(1, 4)
>>> m
array([1, 2, 3])

>>> n = np.arange(4, 10).reshape((2, 3))
>>> n
array([[4, 5, 6],
       [7, 8, 9]])

>>> np.vstack((m, n))
array([[1, 2, 4],
       [4, 5, 6],
       [7, 8, 9]])

```

2.2.6 Splitting

An array split up at N split points will result in N+1 arrays. As for reshaping and concatenation, there are two fundamental ways to split arrays:

1. The function `np.split(array, splitpoints)`.

- `array`: an array of any dimension
- `splitpoints`: a list of indices
 - a divider (positive integer value) can be used to split the array up into n equally sized chunks

2. Functions, which split an array along a specific dimension.

- `np.hsplit(array, splitpoints)`: split the array along the horizontal axis
- `np.vsplit(array, splitpoints)`: split the array along the vertically axis
- `np.dsplit(array, splitpoints)`: split the array along a third dimension

```

>>> a = np.arange(1, 9)
>>> a
array([1, 2, 3, 4, 5, 6, 7, 8])

>>> np.split(a, [4]) # split at index 4 (beginning of second chunk)
[array([1, 2, 3, 4]), array([5, 6, 7, 8])]

```

```

>>> np.split(a, 2) # divide into 2 equally sized parts
[array([1, 2, 3, 4]), array([5, 6, 7, 8])]

>>> np.split(a, [2, 6]) # split at indices 2 and 6
[array([1, 2]), array([3, 4, 5, 6]), array([7, 8])]

>>> b = np.arange(1, 10).reshape((3, 3))
>>> b
array([[1, 2, 3],
       [4, 5, 6],
       [7, 8, 9]])

>>> i, j = np.hsplit(b, [2]) # split off first two columns
>>> i
array([[1, 2],
       [4, 5],
       [7, 8]])

>>> j
array([[3],
       [6],
       [9]])

>>> m, n = np.vsplit(b, [1]) # split off first row
>>> m
array([[1, 2, 3]])

>>> n
array([[4, 5, 6],
       [7, 8, 9]])

```

2.3 Universal Functions

- Loop-based operations on arrays resp. on their elements are slow, because Python performs type-checks and lookups for every function call.
- NumPy's universal functions (UFuncs) are statically typed and compiled. They can be performed on an array as a whole—and will be applied to each element. This is much faster and more convenient.
 - Loops over arrays should be rewritten in terms of UFuncs. The bigger the array, the larger the gain.
- UFuncs can be applied:
 - to an array and a scalar value:
 - * `np.arange(1, 4) * 2` # [2, 4, 6]
 - to an array and another array:
 - * `np.arange(1, 4) * np.arange(7, 10)` # [8, 10, 12]

2.3.1 Common UFuncs

Many of Python's native operators can be used as shorthands for UFuncs:

Shorthand	UFunc	Description
+	np.add	Addition
-	np.subtract	Subtraction
-(unary)	np.negative	Negative Prefix
*	np.multiply	Multiplication
/	np.divide	Division
//	np.floor_divide	Floor Division
**	np.power	Exponentiation
%	np.mod	Modulus (remainder)
np.abs	np.absolute	Absolute value

There are a lot of additional mathematical UFuncs:

- np.sin/np.arcsin: Sine and Arcsine
- np.cos/np.arccos: Cosine and Arcosine
- np.tan/np.arctan: Tangents and Cotangents
- np.exp2: 2^x
- np.exp: e^x
- np.log: base-e logarithm
- np.log2: base-2 logarithm
- np.log10: base-10 logarithm

2.3.2 Advanced Features

Rather than creating a new array for the return value, the result of a UFunc can be stored in an existing array using the out parameter. This also works with slices:

```
>>> x = np.arange(1, 6)
>>> x
array([1, 2, 3, 4, 5])

>>> y = np.zeros(5, dtype=np.int)
>>> y
array([0, 0, 0, 0, 0])

>>> np.power(x, 2, out=y)
>>> y
array([1, 4, 9, 16, 25])

>>> z = np.zeros(10)
>>> z
array([0, 0, 0, 0, 0, 0, 0, 0, 0, 0])

>>> np.power(x, 2, out=z[::2]) # overwrite every other element
>>> z
array([1, 0, 4, 0, 9, 0, 16, 0, 25, 0])
```

Every UFunc comes with a reduce operation, which repeatedly applies an operation to the elements of an array until only a single result remains.

```
>>> x = np.arange(1, 5)
>>> x
array([1, 2, 3, 4, 5])
```

```
>>> np.add.reduce(x) # Sum: 1 + 2 + 3 + 4 + 5
15

>>> np.multiply.reduce(x) # Factorial: 1 * 2 * 3 * 4 * 5
120
```

Instead of just storing the end results, each intermediary step can be stored using the `accumulate` function:

```
>>> x = np.arange(1, 5)
>>> x
array([1, 2, 3, 4, 5])

>>> np.add.accumulate(x)
array([1, 3, 6, 10, 15])

>>> np.multiply.accumulate(x)
array([1, 2, 6, 24, 120])
```

The `outer` operation computes the output of all pairs of two inputs, which could be used to create a multiplication table, for example:

```
>>> a = np.arange(1, 6)
>>> a
array([1, 2, 3, 4, 5])

>>> b = np.arange(1, 9)[1::2]
>>> b
array([2, 4, 6, 8])

>>> np.multiply.outer(b, a) # column, row
array([[ 2,  4,  6,  8, 10],
       [ 4,  8, 12, 16, 20],
       [ 6, 12, 18, 24, 30],
       [ 8, 16, 24, 32, 40]])
```

*	1	2	3	4	5
2	2	4	6	8	10
4	4	8	12	16	20
6	6	12	18	24	30
8	8	16	24	32	40

2.4 Aggregations

Aggregations reduce an array or one of its dimensions to a single value. In contrast to Python's built-in aggregate functions (`sum`, `min`, `max`), NumPy's implementations can operate on multi-dimensional arrays—and are much faster.

- Aggregate functions take an optional `axis` parameter, which describes *the array dimension to be collapsed*:
 - `axis=0`: collapse columns
 - `axis=1`: collapse rows

```

>>> a = np.random.randint(1, 10, size=(3, 4))
>>> a
array([[7, 8, 5, 5],
       [6, 1, 7, 2],
       [7, 2, 8, 8]])

>>> a.sum()
66

>>> a.sum(axis=0)
array([20, 11, 20, 16])

>>> a.sum(axis=1)
array([25, 16, 25])

```

All aggregate functions can be called using the syntax `np.function(array, [parameters])`. Except for `np.median` and `np.percentile`, the following functions can be called directly on the array using the syntax `array.function([parameters])`.

Function	Returns
<code>np.sum</code>	sum
<code>np.prod</code>	product
<code>np.min</code>	minimum value
<code>np.max</code>	maximum value
<code>np.argmin</code>	index of minimum value
<code>np.argmax</code>	index of maximum value
<code>np.mean</code>	mean («average») value
<code>np.median</code>	median («middle») value
<code>np.var</code>	variance
<code>np.std</code>	standard deviation
<code>np.percentile(q=n)</code>	nth percentile, n in [0, 100]
<code>np.any</code>	is <i>any</i> value true?
<code>np.all</code>	are <i>all</i> values true?

Special NaN-aware functions exist for every function (except for the boolean functions `np.any` and `np.all`). They have the prefix `nan` and can only be called on `nd`, not directly on the array. Since NaN belongs to the IEEE-754 standard, arrays containing NaN must have the type `float` or `double`.

```

>>> a = np.array([1, 2, 3, np.NaN, 5])
>>> a
array([ 1.,  2.,  3., nan,  5.])

>>> np.sum(a)
nan

>>> np.nansum(a)
11.0

```

2.5 Broadcasting

Broadcasting is a set of rules for applying binary UFuncs (addition, multiplication, etc.) on arrays of different sizes and/or dimensions.

Rule 1: If the arrays have a different number of dimensions, the *shape* of the array with fewer dimensions is padded with ones on the left.

```
>>> a = np.ones(9).reshape(3, 3)
>>> a
array([[1., 1., 1.],
       [1., 1., 1.],
       [1., 1., 1.]])
```

```
>>> b = np.arange(1, 4)
>>> b
array([1, 2, 3])
```

```
>>> a.shape
(3, 3)
```

```
>>> b.shape
(3,)
```

Result: The shape of b is one-padded on the left: $(3,) \rightarrow (1, 3)$. Thus, `array([1, 2, 3])` becomes `array([[1, 2, 3]])`.

Rule 2: If the shape of the arrays does not match in any dimension, the array with a shape of one is stretched in that dimension to match the other shape.

Result: The rows of b are stretched (i.e. repeated), the shape changes again: $(1, 3) \rightarrow (3, 3)$. Thus, `array([[1, 2, 3]])` becomes:

```
array([[1, 2, 3],
       [1, 2, 3],
       [1, 2, 3]])
```

This is only a *conceptual* transformation, no memory is wasted when stretching!

3. If the dimensions neither match nor are equal to one, an error is raised.

```
>>> x = np.ones(6).reshape(2, 3)
>>> x
array([[1., 1., 1.],
       [1., 1., 1.]])
```

```
>>> y = np.arange(1, 3)
>>> y
array([1, 2])
```

```
>>> x.shape
(2, 3)
```

```
>>> y.shape
(2,)
```

Result: Error.

In order to perform binary operations on incompatible arrays (according these broadcasting rules), the arrays can be re-shaped manually:

```
>>> x + y
ValueError: operands could not be broadcast together with shapes (2,3) (2,)

>>> x + y.reshape(2, 1)
array([[2., 2., 2.],
       [3., 3., 3.]])
```

2.6 Boolean Arrays

Python's comparison operators have NumPy equivalents. They are applied to each element and return a boolean array, indicating the result of every comparison:

Shorthand	UFunc	Description
<code>==</code>	<code>np.equal</code>	equal
<code>!=</code>	<code>np.not_equal</code>	not equal
<code><</code>	<code>np.less</code>	less than
<code>></code>	<code>np.great</code>	greater than
<code><=</code>	<code>np.less_equal</code>	less than or equal
<code>>=</code>	<code>np.greater_equal</code>	greater than or equal

```
>>> a = np.random.randint(1, 10, size=(3, 3))
>>> a
array([[3, 4, 6],
       [7, 4, 2],
       [3, 6, 5]])
```

```
>>> a == 5
array([[False, False, False],
       [False, False, False],
       [False, False,  True]])
```

```
>>> np >= 5
array([[False, False,  True],
       [ True, False, False],
       [False,  True,  True]])
```

```
>>> np.less(a, 5)
array([[ True,  True, False],
       [False,  True,  True],
       [ True, False, False]])
```

The number of true values can be counted using the `np.count_nonzero` or the `np.sum` function, which counts False as 0 and True as 1. Using the array `a` from above:

```
>>> b = a >= 5
>>> b
array([[False, False,  True],
       [ True, False, False],
       [False,  True,  True]])

>>> np.count_nonzero(b)
4
```

```
>>> np.sum(b)
4

>>> np.count_nonzero(b, axis=0)
array([1, 1, 2])

>>> np.sum(b, axis=1)
array([1, 1, 2])
```

2.6.1 Bitmasks

Boolean arrays can be used for indexing, where every True item of the index array is returned:

```
>>> x = np.random.randint(1, 100, size=(4, 4))
>>> x
array([[58, 26, 64, 3],
       [91, 64, 44, 31],
       [14, 81, 77, 8],
       [64, 42, 56, 37]])

>>> above_mean = (x > x.mean())
>>> above_mean
array([[ True, False,  True, False],
       [ True,  True, False, False],
       [False,  True,  True, False],
       [ True, False,  True, False]])

>>> x[above_mean]
array([58, 64, 91, 64, 81, 77, 64, 56])
```

Selection criteria can be combined using the bitwise operands, which are shorthand for NumPy's element-wise logical UFuncs:

Shorthand	UFunc	Description
&	np.bitwise_and	and
\	np.bitwise_or	or
^	np.bitwise_xor	exclusive or
~	np.bitwise_not	not

Using the arrays `x` and `above_mean` from above:

```
>>> even = (x % 2 == 0)
>>> even
array([[ True,  True,  True, False],
       [False,  True,  True, False],
       [ True, False, False,  True],
       [ True,  True,  True, False]])

>>> x[even & above_mean]
array([58, 64, 64, 64, 56])
```

```
>>> x[np.bitwise_or(even, above_mean)]
array([58, 26, 64, 91, 64, 44, 14, 81, 77, 8, 64, 42, 56])

>>> odd = np.bitwise_not(even)
>>> odd
array([[False, False, False,  True],
       [ True, False, False,  True],
       [False,  True,  True, False],
       [False, False, False,  True]])
```

2.7 Fancy Indexing

Arrays can be indexed using arrays of indices to access multiple array elements at once.

```
>>> x = np.arange(5, 85, 5).reshape((4, 4))
>>> x
array([[ 5, 10, 15, 20],
       [25, 30, 35, 40],
       [45, 50, 55, 60],
       [65, 70, 75, 80]])

>>> x[[3, 1, 2], [2, 3, 1]] # select items (3,2), (1,3) and (2,1)
array([75, 40, 50])
```

2.7.1 Broadcasting

If array indices with different shapes are used, the index arrays are being broadcasted. The result of the index operation is shaped by the *broadcasted index array*, not by the array being indexed. Given the array *x* from above:

```
>>> rows = np.array([3, 1, 2])[:, np.newaxis]
>>> rows
array([[3],
       [1],
       [2]])

>>> cols = np.array([2, 3, 1])
>>> cols
array([2, 3, 1])

>>> x[rows, cols]
array([[75, 80, 70],
       [35, 40, 30],
       [55, 60, 50]])
```

The broadcasting of the index arrays is done like this:

	2	3	1
3	3,2	3,3	3,1
1	1,2	1,3	1,1
2	2,2	2,3	2,1

And the resulting array of the indexing operation looks like this:

	2	3	1
3	75	80	70
1	35	40	30
2	55	60	50

Array indices can be combined with scalar indices, slicing and masking:

```
>>> x[2, [1, 0, 3]] # scalar and array index
array([50, 45, 60])

>>> x[2:3, [1, 0, 3]] # slice and array index
array([50, 45, 60])

>>> rows = np.array([2, 3])[:, np.newaxis]
>>> cols = np.array([False, True, False, True])
>>> x[rows, cols] # array index and mask
array([[50, 60],
       [70, 80]])
```

2.7.2 Assignment

Fancy indexing can be used for assignments, too:

```
>>> x = np.arange(10)
>>> x
array([0, 1, 2, 3, 4, 5, 6, 7, 8, 9])

>>> x[x % 2 == 0] = 0 # set all even values to zero
>>> x
array([0, 1, 0, 3, 0, 5, 0, 7, 0, 9])
```

However, the behaviour can be unexpected if index values are used multiple times:

```
>>> x = np.zeros(3)
>>> x
array([0, 0, 0])

>>> i = [0, 1, 1, 2, 2, 2]
>>> x[i] += 1
>>> x
array([1, 1, 1])
```

The values at indices 1 and 2 haven't been incremented three times, because the value of `x[i] + 1` is evaluated once at the beginning and then used multiple times. For repetitions, NumPy's functions have a `at` method, which performs unbuffered operations, i.e. results will be recalculated for every index element:

```
>>> x = np.zeros(3)
>>> x
array([0, 0, 0])

>>> i = [0, 1, 1, 2, 2, 2]
```

```
>>> np.add.at(x, i, 1)
>>> x
array([1, 2, 3])
```

2.8 Sorting

NumPy offers more efficient ways of sorting arrays than Python's native `sort()` function. An array can be sorted using the `np.sort()` function, which returns the sorted array:

```
>>> x = np.array([5, 2, 4, 1, 3])
>>> np.sort(x)
array([1, 2, 3, 4, 5])
```

By default, NumPy uses the quicksort algorithm. Other algorithms can be used by setting the `kind` parameter. Options are: `quicksort`, `mergesort`, `heapsort` and `stable`.

An array can also be sorted in-place, using the array's `sort()` method:

```
>>> x = np.array([5, 2, 4, 1, 3])
>>> x.sort()
>>> x
array([1, 2, 3, 4, 5])
```

The `np.argsort()` function sorts an array and returns an array of indices denoting the array's order. The returned array can be used for fancy indexing:

```
>>> x = np.array([5, 2, 4, 1, 3])
>>> i = np.argsort(x)
>>> i
array([3, 1, 4, 2, 0])

>>> x[i]
array([1, 2, 3, 4, 5])
```

Arrays can be sorted along rows and columns using the `axis` argument, which defines *along (not within!)* which axis the comparison and swapping is performed (0: along rows, 1: along columns):

```
>>> x = np.random.choice(10, (3, 3), replace=False)
>>> x
array([[7, 1, 9],
       [8, 0, 4],
       [2, 3, 6]])

>>> np.sort(x, axis=0) # along rows/within columns
array([[2, 0, 4],
       [7, 1, 6],
       [8, 3, 9]])

>>> np.sort(x, axis=1) # along columns/within rows
array([[1, 7, 9],
       [0, 4, 8],
       [2, 3, 6]])
```

Arrays can be sorted *partially*, i.e. the array is split into two sections, with the left partition containing all smaller values than the right partition. Arrays can be sorted partially using `np.partition()`, which requires the `kth` parameter denoting the size of the left partition (`K` elements):

```
>>> x = np.random.choice(10, 10, replace=False)
>>> x
array([9, 1, 6, 0, 8, 5, 3, 2, 7, 4])

>>> np.partition(x, 3)
array([1, 0, 2, 3, 4, 5, 6, 7, 8, 9])
```

Within the partitions, the elements are in arbitrary order. Partial sorting can also be done by row or column using the `axis` argument. To return the array of partially sorted indices, the function `np.argpartition()` can be used analogous to `np.argsort()`.

2.9 Structured Arrays

Storing heterogeneous data, say names and wages of employees, in different arrays of the same size is error prone: The relation of the data is not obvious, and sorting the arrays mixes up the entries. NumPy offers structured arrays, which can be defined with the `dtype` parameter using a compound data type specification in three ways:

- 1) using the dictionary method, indicating the field names and formats separately in two tuples:

```
dtype={'names': ('name', 'age', 'salary'),
      'formats': ('U20', 'u1', 'f4')}
```

- 2) using a list of tuples, defining the field name and its type together in one tuple per field:

```
dtype=np.dtype([('name', 'U20'),
                ('age', 'u1'),
                ('salary', 'f4')])
```

- 3) without specifying the field names, using automatic names from `f0` to `fn`, and defining the types as a comma-separated string:

```
dtype=np.dtype('U20,u1,f4')
```

A type indicator consists of three parts:

1. the endianness (optional): `<` for little endian, `>` for big endian
 - `<f4`: little endian float of four bytes
 - `>i8`: big endian integer of eight bytes
2. the data type (see the next table)
3. the size of the field *in bytes* (not in bits)

Indicator	Type	Example	Equivalent
'b'	byte	'b'	
'i'	signed integer	'i4'	<code>np.int32</code>
'u'	unsigned integer	'u1'	<code>np.uint8</code>
'f'	floating point	'f8'	<code>np.float64</code>
'c'	complex number	'c16'	<code>np.complex128</code>
'S' or 'a'	string (ASCII)	'S5'	

Indicator	Type	Example	Equivalent
'U'	unicode string	'U10'	np.dtype(np.str_, 10)
'V'	raw data (void)	'V'	np.void

The fields can be accessed by row, by column, by combining row and column, and also using bit masks:

```
>>> employees = np.zeros(3, dtype=np.dtype([('name', 'S10'), ('wage', 'f8')]))
>>> employees['name'] = ['Dilbert', 'Wally', 'Alice']
>>> employees['wage'] = [120000.00, 80000.00, 110000.00]
>>> employees
array([(b'Dilbert', 120000.),
       (b'Wally', 80000.),
       (b'Alice', 110000.)],
      dtype=[('name', 'S10'), ('wage', '<f8')])
```

```
>>> employees[employees['wage'] > 100000]['name']
array([b'Dilbert', b'Alice'], dtype='|S10')
```

NumPy also allows storing arrays in fields of structured arrays, which can be achieved by providing an optional size indicator to every field definition:

```
>>> players = np.zeros(3, dtype=np.dtype([('name', 'U20'),
                                           ('pattern', 'S1', (3, 3))]))
>>> players[0]['name'] = 'John'
>>> glider = [[' ', '0', ' '],
              [' ', ' ', '0'],
              ['0', '0', '0']]
>>> players[0]['pattern'] = glider
```

NumPy offers the type `np.recarray`, which allows the individual fields to be accessed with dot notation instead of array indices:

```
>>> payroll = employees.view(np.recarray)
>>> payroll.name
array([b'Dilbert', b'Wally', b'Alice'], dtype='|S10')
```

The syntax is more convenient, but the performance of the access is lower.

NumPy's structured arrays are a very efficient way to store structured data. However, the Pandas library offers much more functionality for working with structured data.

2.10 Date and Time

Python's capabilities for handling date and time information, such as the modules `datetime`, `dateutil` and `pytz`, are convenient to use, but are too slow when it comes to big datasets.

NumPy defines its own type for that purpose: `datetime64`, which encodes date and time information as 64-bit integers, and can be used for vectorized operations:

```
>>> new_year = np.array('2019-01-01', dtype=np.datetime64)
>>> first_week = new_year + np.arange(7)
>>> first_week
array(['2019-01-01', '2019-01-02', '2019-01-03', '2019-01-04',
       '2019-01-05', '2019-01-06', '2019-01-07'], dtype='datetime64[D]')
```


The `timedelta64` data type is used to express the period between two points in time. Both `datetime64` and `timedelta64` are based on a *fundamental time unit* and can express a range of 2^{64} times that unit. There is a trade-off between resolution (precision) and range (time span): The smaller the fundamental time unit is chosen, the more precision and the less time span can be expressed. The fundamental time unit can be defined as follows:

```
>>> np.datetime64('2019-01-01', 'ns') # 'ns': nanoseconds
numpy.datetime64('2019-01-01T00:00:00.000000000')
```

The options available are:

- Y: year
- M: month
- W: week
- D: day
- h: hour
- m: minute
- s: second
- ms: millisecond
- us: microsecond
- ns: nanosecond
- ps: picosecond
- fs: femtosecond
- as: attosecond

Nanoseconds are a good compromise, for they are as precise as regular computers and have a time span of about 500 years (now ± 250 years).

NumPy infers the time-zone automatically from the operating system.

3 Pandas

Pandas is a package built on top of NumPy, which offers powerful data operations familiar to those of data bases and spreadsheets. The fundamental data structures of Pandas are `Series`, `DataFrame` and `Index`. A `DataFrame` is a multidimensional array with labeled rows and columns, which supports heterogeneous and missing data—an issue often to be faced with in real-world data sets.

Pandas is idiomatically imported as `pd`:

```
>>> import pandas as pd
```

3.1 Series

What NumPy's `ndarray` is to Python's list, Pandas `Series` is to Python's dictionary: a fast and very powerful alternative. Whereas Python's dictionary maps a set of *arbitrary keys* to a set of *arbitrary values*, Pandas `Series` maps a set of *typed keys* to a set of *typed values*. A `Series` is made up of two sequences:

1. values: a NumPy array (`np.ndarray`)
2. index: a Pandas `Index` (`pd.Index`)

3.1.1 Creation

A Pandas Series can be created from scalars, lists and dictionaries.

If a Series is generated from list, the indices (first column) for the values (second column) are made up automatically, i.e. sequentially:

```
>>> pd.Series([1, 2, 3])
0    1
1    2
2    3
dtype: int64
```

An list of indices can be explicitly provided using the index parameter. The the lists of values and indices need to have the same length:

```
>>> pd.Series([1, 2, 3], index=['a', 'b', 'c'])
a    1
b    2
c    3
dtype: int64
```

However, indices can be noncontiguous and nonsequential:

```
>>> pd.Series([1, 2, 3], index=['Foo', 'Bar', 'Qux'])
Foo    1
Bar    2
Qux    3
dtype: int64
```

If a scalar value is used instead of list of values, the same value will be repeated for the length of the index list:

```
>>> pd.Series(42, index=[1, 2, 3])
1    42
2    42
3    42
dtype: int64
```

A Series can be created based on a dictionary with keys to be used as indices:

```
>>> pd.Series({'a': 1, 'b': 2, 'c': 3})
a    1
b    2
c    3
dtype: int64
```

An additional list of indices can be provided to further select values from the dictionary by their keys, and to specify the order of entries:

```
>>> pd.Series({'a': 1, 'b': 2, 'c': 3}, index=['c', 'a'])
c    3
a    1
dtype: int64
```

3.1.2 Access: Indexing and Selection

The elements of a Series can be accessed using indexing and slicing:

```
>>> s = pd.Series([1, 2, 3, 4, 5])
>>> s[0]
1

>>> s[4]
5

>>> s[1:4]
1    2
2    3
3    4
dtype: int64
```

If arbitrary (noncontiguous, nonsequential) indices are used, slicing is possible because of the fixed order of indices, but the upper bound is also included:

```
>>> payroll = pd.Series({'Dilbert': 120000, 'Wally': 80000, 'Alice': 110000})
>>> payroll['Dilbert':'Wally']
Dilbert    120000
Wally       80000
dtype: int64
```

Even though a non-numeric is used, a Series can also be sliced using a implicit index. Here, the upper bound is excluded:

```
>>> payroll[0:2]
Dilbert    120000
Wally       80000
dtype: int64
```

The elements of a Series can also be accessed through the means of masking and fancy indexing:

```
>>> payroll[(payroll >= 100000) & (payroll <= 150000)]
Dilbert    120000
Alice      110000
dtype: int64

>>> payroll[['Alice', 'Dilbert']]
Alice      110000
Dilbert    120000
dtype: int64
```

Python's native dictionary expressions are also supported:

```
>>> 'Dilbert' in payroll
True

>>> 'Asok' in payroll
False

>>> payroll.keys()
Index(['Dilbert', 'Wally', 'Alice'], dtype='object')
```

```
>>> list payroll.items()
[('Dilbert', 120000), ('Wally', 80000), ('Alice', 110000)]

>>> payroll['Wally'] = 90000 # modify existing entry
>>> payroll['Asok'] = 12000 # add a new entry
```

3.1.3 Explicit and Implicit Indexing

When using an explicit integer index, indexing operations make use of the explicit indices (the actual index values provided), but slicing operations use the implicit indices (the items ordinal numbers). This can be confusing:

```
>>> ratings = pd.Series([2.3, 3.1, 3.9, 4.2, 4.8], index=[10, 20, 30, 40, 50])
>>> ratings[10] # explicit index
2.3

>>> ratings[1:3] # implicit index
20    3.1
30    3.9
dtype: float64
```

In order to reduce that confusion, a Series offers two attributes to access the indices:

- loc: the explicit index
- iloc: the implicit index

```
>>> ratings.loc[10]
2.3

>>> ratings.loc[10:30] # inclusive explicit indices from 10 to 30
10    2.3
20    3.1
30    3.9
dtype: float64

>>> ratings.iloc[0]
2.3

>>> ratings.iloc[0:3] # exclusive implicit indices from 0 to 3
10    2.3
20    3.1
30    3.9
dtype: float64
```

According to the [Zen of Python](#) («Explicit is better than implicit.»), slicing and indexing on Series using an integer index should be done using the loc and iloc attributes,

3.2 DataFrame

A Pandas DataFrame can be understood in terms of other data structures from two perspectives:

1. As a generalization of a NumPy array of two dimensions, with row indices and column names being flexible.

- NumPy arrays are indexed as `arr[row, column]`: row first, column second.
- Pandas DataFrames are indexed as `df[column][row]`: column first, row of the Series second.

2. As a specialization of a Python dictionary that maps a column name (key) to a Series of column data (value).

Generally speaking, a DataFrame is a sequence of Series sharing the index value. Important attributes are:

- `columns`: returns an Index object (column names)
- `index`: returns the index labels (row names)

3.2.1 Creation

A Pandas DataFrame can be created from Series, dictionaries and NumPy arrays.

If a single Series is provided, an optional column name for those values can be defined in a list:

```
>>> s = pd.Series([1, 2, 3])
>>> pd.DataFrame(s, columns=['values'])
   values
0        1
1        2
2        3
```

If a list of dictionaries is provided, each dictionary is mapped to a row. Missing entries of heterogeneous dictionaries are filled up with NaN in the resulting DataFrame:

```
>>> pd.DataFrame([{'a': 1, 'b': 2}, {'a': 5, 'c': 4}])
   a  b  c
0  1  2 NaN
1  5 NaN  4
```

If a dictionary of Series is provided, each Series becomes a column with its key mapped as the column name:

```
>>> s1 = pd.Series([2, 4, 6, 8])
>>> s2 = pd.Series([3, 6, 9, 12])
>>> pd.DataFrame({'two': s1, 'three': s2})
   two  three
0     2      3
1     4      6
2     6      9
3     8     12
```

If a two-dimensional NumPy array is provided, the numeric column and row indices from the array are used, but can be set using the optional columns and index parameters:

```
>>> arr = np.arange(1, 10).reshape(3, 3)
>>> pd.DataFrame(arr)
   0  1  2
0  1  2  3
1  4  5  6
2  7  8  9

>>> pd.DataFrame(arr, columns=['A', 'B', 'C'], index=[1, 2, 3])
```

```

      A  B  C
1  1  2  3
2  4  5  6
3  7  8  9

```

If a structured NumPy array is provided, the field names serve as column names:

```

>>> employees = np.zeros(3, dtype=np.dtype([('name', 'S10'), ('wage', 'f8')]))
>>> employees['name'] = ['Dilbert', 'Wally', 'Alice']
>>> employees['wage'] = [120000.00, 80000.00, 110000.00]
>>> pd.DataFrame(employees)
      name      wage
0  b'Dilbert'  120000.0
1    b'Wally'   80000.0
2    b'Alice'  110000.0

```

3.2.2 Access: Indexing and Selection

The DataFrame for the following examples:

```

>>> population = {
...   'USA': 326625792,
...   'Russia': 142257520,
...   'Germany': 80594016,
...   'Switzerland': 8236303
... }

>>> area = {
...   'USA': 9147593,
...   'Russia': 16377742,
...   'Germany': 348672,
...   'Switzerland': 39997
... }

>>> data = pd.DataFrame({'pop': population, 'area': area})
>>> data

```

	pop	area
Germany	80594016	348672
Russia	142257520	16377742
Switzerland	8236303	39997
USA	326625792	9147593

Individual columns can be accessed either dictionary-style or attribute-style, however the latter only works for columns with a string index that isn't used for any other DataFrame attribute:

```

>>> data['area']
Germany      348672
Russia      16377742
Switzerland    39997
USA          9147593
Name: area, dtype: int64

>>> data.area
Germany      348672
Russia      16377742

```

```
Switzerland      39997
USA              9147593
Name: area, dtype: int64
```

```
>>> data['area'] is data.area
True
```

```
>>> data['pop'] is data.pop
False # pop is a method of DataFrame!
```

For assignments, only dictionary-style access works (on the left side):

```
>>> data['density'] = data['pop'] / data.area
>>> data
```

	pop	area	density
Germany	80594016	348672	231.145650
Russia	142257520	16377742	8.686028
Switzerland	8236303	39997	205.923019
USA	326625792	9147593	35.706201

The raw, underlying multi-dimensional array of data of a DataFrame can be accessed using the `values` attribute, which supports array-style indexing:

```
>>> data.values
array([[8.05940160e+07, 3.48672000e+05, 2.31145650e+02],
       [1.42257520e+08, 1.63777420e+07, 8.68602766e+00],
       [8.23630300e+06, 3.99970000e+04, 2.05923019e+02],
       [3.26625792e+08, 9.14759300e+06, 3.57062007e+01]])
```

```
>>> data.values[0, 0]
80594016.0
```

A transposed version of the DataFrame (which rows and columns swapped) can be accessed using the `T` attribute:

```
>>> data.T
```

	Germany	Russia	Switzerland	USA
pop	8.059402e+07	1.422575e+08	8.236303e+06	3.266258e+08
area	3.486720e+05	1.637774e+07	3.999700e+04	9.147593e+06
density	2.311456e+02	8.686028e+00	2.059230e+02	3.570620e+01

A DataFrame offers different index attributes:

- `loc`: explicit index to access values by column and row *names*
 - inclusive upper bound
 - supports name based slicing, masking, fancy indexing
- `iloc`: implicit index to access values by column and row *numbers*
 - zero-based, exclusive upper bound
 - supports row and column access by ordinal numbers

```
>>> data.loc['Germany':'Russia', 'pop':'area']
```

	pop	area
Germany	80594016	348672
Russia	142257520	16377742

```
>>> data.loc[data.density > 100, ['pop', 'density']]
```

	pop	density
--	-----	---------

```
Germany      80594016  231.145650
Switzerland  8236303   205.923019
```

```
>>> data.iloc[0:2, 0:2]
      pop      area
Germany  80594016  348672
Russia  142257520  16377742
```

3.3 Index

The Pandas Index is an immutable array/a ordered (multi)set that is used both for the indexing of Series and DataFrame.

An Index can be created from a list:

```
>>> pd.Index([1, 2, 3, 4, 5])
Int64Index([1, 2, 3, 4, 5], dtype='int64')
```

The elements of the Index can be accessed like list entries, i.e. by a single index and using slicing:

```
>>> idx = pd.Index([1, 2, 3, 4, 5])
>>> idx[2]
3
```

```
>>> idx[0:2]
Int64Index([1, 2], dtype='int64')
```

```
>>> idx[::2]
Int64Index([1, 3, 5], dtype='int64')
```

An Index is immutable, which is important when they are shared between different DataFrames and Series:

```
>>> idx[2] = 6
TypeError: Index does not support mutable operations
```

Like Python's native set, Index supports set operations like intersection, union and difference:

```
>>> idxA.intersection(idxB)
Int64Index([1, 3, 5], dtype='int64')
```

```
>>> idxA.union(idxB)
Int64Index([1, 2, 3, 4, 5, 7, 9], dtype='int64')
```

```
>>> idxA.difference(idxB)
Int64Index([7, 9], dtype='int64')
```

```
>>> idxB.difference(idxA)
Int64Index([2, 4], dtype='int64')
```

```
>>> idxA.symmetric_difference(idxB)
Int64Index([2, 4, 7, 9], dtype='int64')
```

Union, intersection and symmetric difference can be expressed by the means of operators:


```
>>> idxA & idxB # intersection
Int64Index([1, 3, 5], dtype='int64')

>>> idxA | idxB # union
Int64Index([1, 2, 3, 4, 5, 7, 9], dtype='int64')

>>> idxA ^ idxB # symmetric difference
Int64Index([2, 4, 7, 9], dtype='int64')
```

3.4 Operations

Pandas offers a lot of functions like NumPy's UFuncs that can be applied on a Series or DataFrame either using a method (with another Series or DataFrame as a argument) or using a Python operator:

Operator	Method	Description
+	<code>add()</code>	Addition
-	<code>sub()</code> , <code>subtract()</code>	Subtraction
*	<code>mul()</code> , <code>multiply()</code>	Multiplication
/	<code>truediv()</code> , <code>div()</code> , <code>divide()</code>	Division
//	<code>floordiv()</code>	Floor Division
%	<code>mod()</code>	Modulus (remainder)
**	<code>pow()</code>	Exponentiation

The index of the operands is preserved in the result. If the operands are heterogeneous, the result contains the union of the two indices, with NaN filled in for missing values:

```
>>> hours = pd.Series([25, 40, 32], index=['Alice', 'Bob', 'Malory'])
>>> rates = pd.Series([45, 50, 30], index=['Alice', 'Bob', 'Thomas'])
>>> hours * rates
Alice    1125.0
Bob      2000.0
Malory      NaN
Thomas      NaN
dtype: float64
```

An operation that mixes a Series and a DataFrame works like an operation on a one-dimensional and a multi-dimensional array; broadcasting rules (similar as those for NumPy) apply:

```
>>> wages = pd.DataFrame({'January': {'Alice': 4500, 'Bob': 4800},
...                        'February': {'Alice': 4200, 'Bob': 4500}})
>>> wages
      January  February
Alice    4500     4200
Bob      4800     4500

>>> increase = pd.Series({'Alice': 1.2, 'Bob': 1.1})
>>> increase
Alice    1.2
Bob      1.1
dtype: float64
```

```
>>> wages.T * increase # with transposition
      Alice    Bob
January 5400.0 5280.0
February 5040.0 4950.0

>>> wages.multiply(increase, axis=0) # with optional axis (increase as rows)
      January  February
Alice  5400.0   5040.0
Bob    5280.0   4950.0
```

Pandas always preserves indices and column names, so that the data context is maintained.

3.5 Handling Missing Data

Real-world data sets are rarely clean and homogeneous. Oftentimes, values are missing, and the lack of a value is indicated in different ways. Pandas marks the absence of a value in two different ways:

1. `None`: a Python singleton object, which is used in object collections (rather slow due to the overhead).
2. `NaN`: a special floating point value (not a number), which is defined in the IEEE-754 standard and used for numeric collections. NumPy's `NaN` reference is used: `np.nan`.

A `Series` and `DataFrame` containing a `None` or `NaN` «value» is upcast according to the types of the other elements: integer types are upcast to `float64`; booleans are upcast to `object`.

```
>>> pd.Series([1, 2, None]) # None replaced by NaN
0    1.0
1    2.0
2    NaN
dtype: float64

>>> pd.Series([1, 2, np.nan])
0    1.0
1    2.0
2    NaN
dtype: float64

>>> pd.Series([True, False, None]) # None preserved
0    True
1   False
2    None
dtype: object

>>> pd.Series([True, False, np.nan])
0    True
1   False
2    NaN
dtype: object
```

Any operation involving `NaN` yields `NaN`:

```
>>> 3 + np.nan
nan
```

```
>>> (3 + 7) * np.nan
nan
```

```
>>> pd.Series([1, 2, np.nan]) + pd.Series([1, np.nan, 3])
0    2.0
1    NaN
2    NaN
dtype: float64
```

Whereas NumPy supports special NaN-aware functions (`np.nansum()`, `np.nanmax()`), Pandas offers special functions to deal with absent values:

`isnull()` and `notnull()` return a boolean mask indicating if there is no value (`isnull`) or a value (`notnull`) at the respective index. These masks can be used for indexing:

```
>>> s = pd.Series([1, np.nan, 3])
>>> s.isnull()
0    False
1     True
2    False
dtype: bool
```

```
>>> s.notnull()
0     True
1    False
2     True
dtype: bool
```

```
>>> s[s.notnull()]
0    1.0
2    3.0
dtype: float64
```

`dropna()` removes None and NaN entries in a Series. In a DataFrame, the full row or column missing a value is removed, which can be defined using the optional `axis` parameter:

```
>>> farmers = ['Miller', 'Shaw', 'Watson']
>>> dogs = pd.Series([1, 2, 1], index=farmers)
>>> cats = pd.Series([3, 1, np.nan], index=farmers)
>>> cows = pd.Series([7, np.nan, 2], index=farmers)
>>> pigs = pd.Series([0, 2, np.nan], index=farmers)
>>> livestock = pd.DataFrame( {'dogs': dogs, 'cats': cats, 'cows': cows, 'pigs': pigs})
>>> livestock
```

	dogs	cats	cows	pigs
Miller	1	3.0	7.0	0.0
Shaw	2	1.0	NaN	2.0
Watson	1	NaN	2.0	NaN

```
>>> livestock.dropna() # default: axis='rows'
      dogs  cats  cows  pigs
Miller    1   3.0   7.0   0.0
```

```
>>> livestock.dropna(axis='columns')
      dogs
Miller    1
Shaw      2
Watson    1
```

By default, every row/column with at least one missing entry is dropped. If the optional how parameter is set to all, only rows/columns with missing values only are dropped:

```
>>> livestock.dropna() # default: how='any'
```

	dogs	cats	cows	pigs
Miller	1	3.0	7.0	0.0

```
>>> livestock.dropna(how='all')
```

	dogs	cats	cows	pigs
Miller	1	3.0	7.0	0.0
Shaw	2	1.0	NaN	2.0
Watson	1	NaN	2.0	NaN

The optional parameter thresh allows to define a threshold: only drop rows/columns with fewer values given:

```
>>> livestock.dropna(thresh=3) # drop rows with fewer than three values
```

	dogs	cats	cows	pigs
Miller	1	3.0	7.0	0.0
Shaw	2	1.0	NaN	2.0

```
>>> livestock.dropna(thresh=3, axis='columns')
```

	dogs
Miller	1
Shaw	2
Watson	1

fillna() fills in a value where one is missing. Either a scalar value can be passed, or the value from a neighbouring cell can be propagated using a combination of the method (ffill/bfill: forward and backward fill) and axis (rows/columns) parameters:

```
>>> livestock.fillna(0) # replace NaN with 0, which is useful for sums
```

	dogs	cats	cows	pigs
Miller	1	3.0	7.0	0.0
Shaw	2	1.0	0.0	2.0
Watson	1	0.0	2.0	0.0

```
>>> livestock.fillna(method='ffill', axis='rows') # propagate value to next row
```

	dogs	cats	cows	pigs
Miller	1	3.0	7.0	0.0
Shaw	2	1.0	7.0	2.0
Watson	1	1.0	2.0	2.0

```
>>> livestock.fillna(method='bfill', axis='columns') # ... from previous column
```

	dogs	cats	cows	pigs
Miller	1.0	3.0	7.0	0.0
Shaw	2.0	1.0	2.0	2.0
Watson	1.0	2.0	2.0	NaN

If there is no next or previous row or column, NaN entries could still remain after the fillna() operation.

3.6 Hierarchical Indexing

Pandas Series and DataFrame represent one- and two-dimensional data. But some data must be indexed by more than two indices, and values can only be accessed by a combination of all those indices. This concept is called *hierarchical indexing* or *multi-indexing*.

A index with multiple levels could be represented by a tuple (using Formula 1 teams and seasons as indices):

```
>>> index = [
    ('Mercedes', 2018), ('Mercedes', 2017),
    ('Ferrari', 2018), ('Ferrari', 2017),
    ('McLaren', 2018), ('McLaren', 2017)]
>>> points = pd.Series([655, 688, 571, 522, 62, 30], index=index)
>>> points
(Mercedes, 2018)    655
(Mercedes, 2017)    688
(Ferrari, 2018)     571
(Ferrari, 2017)     522
(McLaren, 2018)      62
(McLaren, 2017)     30
dtype: int64
```

However, storing a tuple as the index is inconvenient and inefficient for data access. Therefore Pandas offers MultiIndex, an efficient wrapper for tuple indices:

```
>>> multi_index = pd.MultiIndex.from_tuples(index)
>>> multi_index
MultiIndex(levels=[['Ferrari', 'McLaren', 'Mercedes'], [2017, 2018]],
            labels=[[2, 2, 0, 0, 1, 1], [1, 0, 1, 0, 1, 0]])
```

The MultiIndex has two levels (the team names and seasons), and they are combined with labels like this:

Team	labels[0]	labels[1]	Season
Mercedes	2	1	2018
Mercedes	2	0	2017
Ferrari	0	1	2018
Ferrari	0	0	2017
McLaren	1	1	2018
McLaren	1	0	2017

A Series created with a tuple index can be reindexed using a MultiIndex:

```
>>> points = points.reindex(multi_index)
>>> points
Mercedes 2018    655
          2017    688
Ferrari   2018    571
          2017    522
McLaren   2018     62
          2017     30
dtype: int64
```

The blank space below the team index means that the value from above is used.

A DataFrame with additional columns can be created based on the existing DataFrame:

```
>>> f1 = pd.DataFrame({
    'points': points,
    'races': [21, 20, 21, 20, 21, 20],
    'wins': [11, 12, 6, 5, 0, 0]})
>>> f1
```

		points	races	wins
Mercedes	2018	655	21	11
	2017	688	20	12
Ferrari	2018	571	21	6
	2017	522	20	5
McLaren	2018	62	21	0
	2017	30	20	0

The operations mentioned earlier can also be applied:

```
>>> win_ratio = f1['wins'] / f1['races']
>>> win_ratio
```

Mercedes	2018	0.523810
	2017	0.600000
Ferrari	2018	0.285714
	2017	0.250000
McLaren	2018	0.000000
	2017	0.000000

dtype: float64

3.6.1 Creation of Hierarchical Indices

A hierarchical index can be created implicitly, i.e. together with the Series or the DataFrame.

The index can be passed as an additional argument to the constructor as a list of index arrays:

```
>>> points = [655, 688, 571, 522]
>>> index = [['Mercedes', 'Mercedes', 'Ferrari', 'Ferrari'],
             [2018, 2017, 2018, 2017]]
>>> pd.Series(points, index=index)
```

Mercedes	2018	655
	2017	688
Ferrari	2018	571
	2017	522

dtype: int64

Or a dictionary can be passed to the constructor, with appropriate index tuples as keys:

```
>>> points = {('Mercedes', 2018): 655, ('Mercedes', 2017): 688,
              ('Ferrari', 2018): 571, ('Ferrari', 2017): 522}
>>> pd.Series(points)
```

Mercedes	2018	655
	2017	688
Ferrari	2018	571
	2017	522

dtype: int64

Using one of MultiIndex class methods, a hierarchical index can be created explicitly. The resulting object can be passed to the constructor of a Series or a DataFrame as the index attribute.

The method `from_arrays` accepts a list of index arrays:

```
>>> pd.MultiIndex.from_arrays([['Mercedes', 'Mercedes', 'Ferrari', 'Ferrari'],
                               [2018, 2017, 2018, 2017]])
MultiIndex(levels=[['Ferrari', 'Mercedes'], [2017, 2018]],
            labels=[[1, 1, 0, 0], [1, 0, 1, 0]])
```

The method `from_tuples` accepts a list of index tuples:

```
>>> pd.MultiIndex.from_tuples([('Mercedes', 2018), ('Mercedes', 2017),
                               ('Ferrari', 2018), ('Ferrari', 2017)])
MultiIndex(levels=[['Ferrari', 'Mercedes'], [2017, 2018]],
            labels=[[1, 1, 0, 0], [1, 0, 1, 0]])
```

In the above examples, every item from the first index (['Mercedes', 'Ferrari']) has been combined with every item from the second index ([2018, 2017]) *manually*. This Cartesian product can also be created automatically using the `from_product` method:

```
>>> pd.MultiIndex.from_product([['Mercedes', 'Ferrari'], [2018, 2017]])
MultiIndex(levels=[['Ferrari', 'Mercedes'], [2017, 2018]],
            labels=[[1, 1, 0, 0], [1, 0, 1, 0]])
```

The index levels can also be combined manually using a nested list of labels passed to the constructor of MultiIndex. This is especially helpful, if only certain combinations of index entries need to be created:

```
>>> index = pd.MultiIndex(levels=[['Manor', 'Haas'], [2015, 2016, 2017]],
                           labels=[[0,0,1,1], [0,1,1,2]])
>>> pd.Series([0, 1, 29, 47], index=index)
Manor  2015    0
       2016    1
Haas   2016   29
       2017   47
dtype: int64
```

Using a DataFrame, both rows and columns can have multiple indices:

```
>>> row_index = pd.MultiIndex.from_product([['Mercedes', 'Ferrari'],
                                             [2018, 2017]])
>>> col_index = pd.MultiIndex.from_product([['Australia', 'Bahrain'],
                                             ['Driver 1', 'Driver 2']])
>>> pos = np.array([2, 8, 3, 2, 2, 3, 2, 3, 1, 3, 1, np.nan, 1, 4, 1, 4])
>>> f1 = pd.DataFrame(pos.reshape((4, 4)), index=row_index, columns=col_index)
>>> f1
```

		Australia		Bahrain	
		Driver 1	Driver 2	Driver 1	Driver 2
Mercedes	2018	2.0	8.0	3.0	2.0
	2017	2.0	3.0	2.0	3.0
Ferrari	2018	1.0	3.0	1.0	NaN
	2017	1.0	4.0	1.0	4.0

This allows for four-dimensional indices.

Both row and column index can be named by setting a list of row/column names with the appropriate length to the `names` attribute of the index:

```
>>> f1.index.names = ['Team', 'Season']
>>> f1.columns.names = ['GP', 'Driver']
>>> f1
```

		Australia		Bahrain	
		Driver 1	Driver 2	Driver 1	Driver 2
Team	Season				
Mercedes	2018	2.0	8.0	3.0	2.0
	2017	2.0	3.0	2.0	3.0
Ferrari	2018	1.0	3.0	1.0	NaN
	2017	1.0	4.0	1.0	4.0

3.6.2 Indexing and Slicing

Indexing and Slicing on Series is row based. This Series index has a species as the first (higher level) index, and the year as the second (lower level) index:

```
>>> idx = pd.MultiIndex.from_product(['cats', 'cows', 'dogs', 'pigs'],
                                     [2000, 2005, 2010])
>>> livestock = pd.Series([32, 16, 25, 60, 75, 52, 1, 1, 2, 4, 3, 7], index=idx)
>>> livestock
```

cows	2000	32
	2005	16
	2010	25
pigs	2000	60
	2005	75
	2010	52
dogs	2000	1
	2005	1
	2010	2
cats	2000	4
	2005	3
	2010	7

```
dtype: int64
```

Individual values can be accessed using full indexing by first indicating the higher level index and second the lower level index:

```
>>> livestock['cats', 2000]
4
>>> livestock['cows', 2010]
25
>>> livestock['pigs', 2005] - livestock['pigs', 2010]
23
```

If the lower level index is left unspecified, a Series with the lower level index retained is returned:

```
>>> livestock['cows']
```

2000	32
2005	16
2010	25

```
dtype: int64
```

Passing an empty slice for the higher level index allows indexing on the lower level index:


```
>>> livestock[:, 2010]
cows      25
pigs      52
dogs       2
cats       7
dtype: int64
```

Slicing on the explicit index is only available on a dataset with a sorted MultiIndex. Either the dataset is created using a sorted MultiIndex:

```
>>> idx = idx.sort_values()
>>> livestock = pd.Series([4, 3, 7, 32, 16, 25, 1, 1, 2, 60, 75, 52], index=idx)
```

Or the MultiIndex on the existing dataset is sorted, returning a new dataset:

```
>>> livestock = livestock.sort_index()
```

The indices are sorted lexicographically. Then the slicing operations can be performed (on the explicit index):

```
>>> livestock.loc['cats':'cows', 2000:2005]
cats  2000      4
      2005      3
cows  2000     32
      2005     16
dtype: int64
```

Selections can be made based on boolean masks:

```
>>> livestock[livestock > 10]
cows  2000     32
      2005     16
      2010     25
pigs  2000     60
      2005     75
      2010     52
dtype: int64
```

Values can be selected using fancy indexing:

```
>>> livestock[['cows', 'pigs']]
cows  2000     32
      2005     16
      2010     25
pigs  2000     60
      2005     75
      2010     52
dtype: int64
```

The indexing hierarchy on a DataFrame behaves like the one of a Series, expect that a DataFrame is indexed by columns first:

```
>>> row_idx = pd.MultiIndex.from_product([[2017, 2018],
                                           ['Jan', 'Jul']])
>>> col_idx = col_idx = pd.MultiIndex.from_product(['Tom', 'Jim'],
                                                    ['height', 'weight'])
>>> val = [[122, 35, 129, 37],
           [128, 37, 131, 39],
           [134, 39, 135, 41],
```

```

[137, 40, 138, 43]]
>>> kids = pd.DataFrame(val, columns=col_idx, index=row_idx)
>>> kids
      Tom      Jim
      height weight height weight
2017 Jan   122    35   129    37
      Jul   128    37   131    39
2018 Jan   134    39   135    41
      Jul   137    40   138    43

>>> kids['Tom', 'height']
2017 Jan   122
      Jul   128
2018 Jan   134
      Jul   137
Name: (Tom, height), dtype: int64

```

For row-oriented selection on a DataFrame, the implicit index can be used:

```

>>> kids.iloc[0:2]
      Tom      Jim
      height weight height weight
2017 Jan   122    35   129    37
      Jul   128    37   131    39

```

The column index hierarchy can be expressed using the explicit index and tuples:

```

>>> kids.loc[:, ('Tom', 'weight')]
2017 Jan    35
      Jul    37
2018 Jan    39
      Jul    40
Name: (Tom, weight), dtype: int64

```

Because tuples do not support slices, Pandas offers the IndexSlice object:

```

>>> jan = pd.IndexSlice[:, 'Jan']
>>> weight = pd.IndexSlice[:, 'weight']
>>> kids.loc[jan, weight]
      Tom      Jim
      weight weight
2017 Jan    35    37
2018 Jan    39    41

```

3.6.3 Rearranging Multi-Indices

Conceptually, a Series with two indices is a lot like a DataFrame, which maps the first index to the rows and the second index to the columns. A multi-index Series can be converted to a DataFrame using the Series `unstack()` method:

```

>>> idx = pd.MultiIndex.from_product([[2017, 2018],
                                      ['Bezos', 'Gates', 'Buffet']])
>>> billions = [72.8, 75.6, 86.0, 112, 84, 90]
>>> richest = pd.Series(billions, index=idx.sort_values())
>>> richest
2017 Bezos    72.8

```

```

    Buffet    75.6
    Gates     86.0
2018 Bezos    112.0
    Buffet    84.0
    Gates     90.0
dtype: float64

```

```

>>> richest.unstack()
      Bezos  Buffet  Gates
2017   72.8   75.6   86.0
2018  112.0   84.0   90.0

```

An optional level can be defined to indicate which index level is to be transformed into a column level:

```

>>> richest.unstack(level=0)
      2017  2018
Bezos   72.8 112.0
Buffet   75.6 84.0
Gates    86.0 90.0

```

```

>>> richest.unstack(level=1)
      Bezos  Buffet  Gates
2017   72.8   75.6   86.0
2018  112.0   84.0   90.0

```

The DataFrame can be converted back to a multi-index Series using the `stack()` method. The column index will become the lower level index of the row MultiIndex:

```

>>> richest.unstack(level=0).stack()
Bezos    2017    72.8
        2018   112.0
Buffet   2017    75.6
        2018    84.0
Gates    2017    86.0
        2018    90.0
dtype: float64

```

The indices of a dataset can be turned into regular columns using the `reset_index()` method, which allows to name the existing data column using an optional argument:

```

>>> richest.index.names = ['year', 'person']
>>> table = richest.reset_index(name='billions')
>>> table
   year  person  billions
0  2017   Bezos    72.8
1  2017  Buffet    75.6
2  2017   Gates    86.0
3  2018   Bezos   112.0
4  2018  Buffet    84.0
5  2018   Gates    90.0

```

Data columns can also be turned (back) into a MultiIndex using the `set_index()` method, which expects a list of columns to be used as indices:

```

>>> table.set_index(['year', 'person'])
billions
year person

```

2017	Bezos	72.8
	Buffet	75.6
	Gates	86.0
2018	Bezos	112.0
	Buffet	84.0
	Gates	90.0

Aggregation methods have optional `level` and `axis` parameters, which allow for partial aggregations:

```
>>> richest.mean(level='year')
year
2017    78.133333
2018    95.333333
dtype: float64
```

```
>>> richest.mean(level='person')
person
Bezos    92.4
Buffet   79.8
Gates    88.0
dtype: float64
```

```
>>> richest.unstack(level=0).mean(axis=0)
year
2017    78.133333
2018    95.333333
dtype: float64
```

```
>>> richest.unstack(level=0).mean(axis=1)
person
Bezos    92.4
Buffet   79.8
Gates    88.0
dtype: float64
```

`level` and `axis` can also be combined, which is useful if both row and column use a `MultiIndex`.

3.6.4 Multi-Indices vs. Panels

Datasets using a `MultiIndex` are *sparse representations* of data: only the existing values are represented. Panels (classes `Panel` and `Panel4D`), in contrast, are *dense representations* of data. A value is stored for every combination of all indices. Since real-world data sets are often sparse, `MultiIndex` datasets are often more efficient than panels.

3.7 Combining Datasets

Conducting interesting studies of data often requires combining datasets from different sources. Pandas offers different facilities to perform this task: concatenations and database-style joins.

3.7.1 Concat and Append

To demonstrate the concatenation of datasets, this function is used to create a DataFrame quickly with values made up of column names and row indices:

```
def create_df(cols, index):
    data = {c: [str(c) + str(i) for i in index] for c in cols}
    return pd.DataFrame(data, index)
```

The function can be used thus:

```
>>> create_df('ABC', range(3))
```

	A	B	C
0	A0	B0	C0
1	A1	B1	C1
2	A2	B2	C2

Multiple Series or DataFrames can be combined using Pandas concat function, which expects a list of datasets:

```
>>> a = create_df('ABC', [1, 2, 3])
>>> b = create_df('ABC', [4, 5, 6])
>>> pd.concat([a, b])
```

	A	B	C
1	A1	B1	C1
2	A2	B2	C2
3	A3	B3	C3
4	A4	B4	C4
5	A5	B5	C5
6	A6	B6	C6

By default, the concatenation is performed row-wise (default parameter axis=0). The concatenation can be performed column-wise by setting the axis parameter either to 1:

```
>>> a = create_df('ABC', [1, 2, 3])
>>> b = create_df('DEF', [1, 2, 3])
>>> pd.concat([a, b], axis=1)
```

	A	B	C	D	E	F
1	A1	B1	C1	D1	E1	F1
2	A2	B2	C2	D2	E2	F2
3	A3	B3	C3	D3	E3	F3

By default, indices are preserved, even if the resulting index contains duplicates:

```
>>> a = create_df('ABC', [0, 1, 2])
>>> b = create_df('ABC', [2, 3, 4])
>>> pd.concat([a, b])
```

	A	B	C
0	A0	B0	C0
1	A1	B1	C1
2	A2	B2	C2
2	A2	B2	C2
3	A3	B3	C3
4	A4	B4	C4

The index 2 occurs twice in the resulting dataset above. There are different ways to deal with duplicate indices. The first is to raise an error in case of conflict by setting the verify_integrity flag to True:

```
>>> pd.concat([a, b], verify_integrity=True)
ValueError: Indexes have overlapping values: Int64Index([2], dtype='int64')
```

An other option is to ignore the existing indices and let Pandas create a new one by setting the `ignore_index` flag to `True`:

```
>>> pd.concat([a, b], ignore_index=True)
   A  B  C
0  A0 B0 C0
1  A1 B1 C1
2  A2 B2 C2
3  A2 B2 C2
4  A3 B3 C3
5  A4 B4 C4
```

The existing indices can be converted to a `MultiIndex` by introducing a higher-level index key describing the source of the entries in the resulting dataset using the `keys` parameter:

```
>>> pd.concat([a, b], keys=['a', 'b'])
      A  B  C
a 0  A0 B0 C0
  1  A1 B1 C1
  2  A2 B2 C2
b 2  A2 B2 C2
  3  A3 B3 C3
  4  A4 B4 C4
```

If datasets with columns in common are concatenated, the resulting dataset is a union of the source datasets (default parameter `join='outer'`). Missing values (in uncommon columns) are filled up as `NaN`:

```
>>> a = create_df('ABC', range(3))
>>> b = create_df('BCD', range(3))
>>> pd.concat([a, b])
      A  B  C  D
0  A0  B0 C0 NaN
1  A1  B1 C1 NaN
2  A2  B2 C2 NaN
0  NaN  B0 C0 D0
1  NaN  B1 C1 D1
2  NaN  B2 C2 D2
```

If the resulting dataset should only consist of the columns in common of the source datasets, setting the parameter `join='inner'` will create a dataset as an intersection of the source columns:

```
>>> pd.concat([a, b], join='inner')
      B  C
0  B0  C0
1  B1  C1
2  B2  C2
0  B0  C0
1  B1  C1
2  B2  C2
```

For fine-grained control of the resulting columns, the parameter `join_axes` can be set to a `Index` object representing the output columns:

```
>>> pd.concat([a, b], join_axes=[pd.Index(['A', 'B', 'C'])])
   A  B  C
0  A0 B0 C0
1  A1 B1 C1
2  A2 B2 C2
0  NaN B0 C0
1  NaN B1 C1
2  NaN B2 C2
```

An existing Index object of the source datasets can also be used:

```
>>> pd.concat([a, b], join_axes=a.columns)
   A  B  C
0  A0 B0 C0
1  A1 B1 C1
2  A2 B2 C2
0  NaN B0 C0
1  NaN B1 C1
2  NaN B2 C2
```

The `append()` method of a DataFrame is a shorthand for the `pd.concat()` function:

```
>>> a = create_df('ABC', range(3))
>>> b = create_df('ABC', [3, 4, 5])
>>> a.append(b)
   A  B  C
0  A0 B0 C0
1  A1 B1 C1
2  A2 B2 C2
3  A3 B3 C3
4  A4 B4 C4
5  A5 B5 C5
```

It should not be used when combining more than two datasets, because new indices and data buffers are created for every intermediary step.

3.7.2 Merge and Join

Pandas offers high-performance, in-memory join and merge operations. The `pd.merge()` function is the main interface, but `DataFrame` and `Series` also offer a `join()` method for higher convenience.

There are three types of joins:

1. one-to-one (1:1)
2. one-to-many (1:n)
3. many-to-many (n:m)

The type of join to be performed depends solely on the input data.

A one-to-one join is similar to column-wise concatenation. The datasets are automatically joined using a column common to both datasets:

```
>>> employees = pd.DataFrame(
    {'employee': ['Dilbert', 'Catbert', 'Pointy Haired Boss'],
     'department': ['Engineering', 'HR', 'Management']})
>>> employees
   employee  department
```

```

0          Dilbert  Engineering
1          Catbert           HR
2 Pointy Haired Boss  Management

>>> departments = pd.DataFrame(
    {'department': ['Management', 'HR', 'Engineering'],
     'location': ['upper floor', 'middle floor', 'basement']})
>>> departments
   department  location
0  Management  upper floor
1         HR  middle floor
2  Engineering    basement

>>> pd.merge(employees, departments)
   employee  department  location
0     Dilbert  Engineering    basement
1     Catbert           HR  middle floor
2 Pointy Haired Boss  Management  upper floor

```

The index of the input datasets is discarded; a new index is generated for the resulting dataset. The order of entries in the output may be different from the input.

If one of the key columns contains duplicates, a one-to-many join is performed. Using the same departments, but a extended employees dataset:

```

>>> employees = pd.DataFrame(
    {'employee': ['Dilbert', 'Wally', 'Catbert', 'Pointy Haired Boss'],
     'department': ['Engineering', 'Engineering', 'HR', 'Management']})
>>> employees
   employee  department
0     Dilbert  Engineering
1       Wally  Engineering
2     Catbert           HR
3 Pointy Haired Boss  Management

>>> pd.merge(employees, departments)
   employee  department  location
0     Dilbert  Engineering    basement
1       Wally  Engineering    basement
2     Catbert           HR  middle floor
3 Pointy Haired Boss  Management  upper floor

```

If the key columns on both sides contain duplicates, a many-to-many join is performed:

```

>>> employees = pd.DataFrame(
    {'name': ['Dilbert', 'Wally', 'Catbert'],
     'department': ['Engineering', 'Engineering', 'HR']})
>>> employees
   name  department
0  Dilbert  Engineering
1   Wally  Engineering
2  Catbert           HR

>>> skills = pd.DataFrame(
    {'skill': ['programming', 'thinking', 'thinking', 'manipulating'],
     'department': ['Engineering', 'Engineering', 'HR', 'HR']})
>>> skills

```


	skill	department
0	programming	Engineering
1	thinking	Engineering
2	thinking	HR
3	manipulating	HR

```
>>> pd.merge(employees, skills)
   name  department      skill
0  Dilbert  Engineering  programming
1  Dilbert  Engineering    thinking
2    Wally  Engineering  programming
3    Wally  Engineering    thinking
4  Catbert         HR      thinking
5  Catbert         HR  manipulating
```

These examples all assume *one column common to both datasets*, which is often not given in real-world datasets. The behaviour of `merge()` can be further specified to overcome this constraint.

If there are multiple common columns in both datasets, the column to be joined on can be defined using the `on` parameter:

```
>>> employees = pd.DataFrame(
    {'id': [1, 2, 3],
     'name': ['Dilbert', 'Wally', 'Catbert'],
     'department': ['Engineering', 'Engineering', 'HR']})
>>> employees
```

	id	name	department
0	1	Dilbert	Engineering
1	2	Wally	Engineering
2	3	Catbert	HR

```
>>> departments = pd.DataFrame(
    {'id': [1, 2],
     'department': ['Engineering', 'HR'],
     'location': ['basement', 'middle floor']})
>>> departments
```

	id	department	location
0	1	Engineering	basement
1	2	HR	middle floor

```
>>> pd.merge(employees, departments, on='department')
   id_x  name  department  id_y  location
0     1  Dilbert  Engineering     1  basement
1     2    Wally  Engineering     1  basement
2     3  Catbert         HR      2  middle floor
```

If the columns to be joined have a different name, the join can be defined using the `left_on` and `right_on` parameters:

```
>>> employees = pd.DataFrame(
    {'id': [1, 2, 3],
     'name': ['Dilbert', 'Wally', 'Catbert'],
     'department_id': [1, 1, 2]})
>>> employees
```

	id	name	department_id
0	1	Dilbert	1

```

1  2  Wally          1
2  3  Catbert        2

>>> departments = pd.DataFrame(
    {'id': [1, 2, 3],
     'department': ['Engineering', 'HR', 'Management']})
>>> departments
   id  department
0   1  Engineering
1   2           HR
2   3  Management

>>> pd.merge(employees, departments,
             left_on='department_id', right_on='id')
   id_x  name  department_id  id_y  department
0     1  Dilbert           1     1  Engineering
1     2   Wally           1     1  Engineering
2     3  Catbert           2     2           HR

```

Redundant columns can be removed from the output using the `drop()` method by providing the name of the column to be discarded, and the argument `axis=1` to specify that the column has to be dropped (as opposed to the row with `axis=0`):

```

>>> pd.merge(employees, departments,
             left_on='department_id', right_on='id').drop('id_x', axis=1)
   name  department_id  id_y  department
0  Dilbert           1     1  Engineering
1   Wally           1     1  Engineering
2  Catbert           2     2           HR

```

Joins can also be performed based on the index instead of on columns. Using the datasets `employees` and `department` from above with appropriate indices, the join can be performed by setting the `left_index` and `right_index` flags to `True`:

```

>>> employees = employees.set_index('id')
>>> employees
      name  department_id
id
1  Dilbert           1
2   Wally           1
3  Catbert           2

>>> departments = departments.set_index('id')
>>> departments
      department
id
1  Engineering
2           HR
3  Management

>>> pd.merge(employees, departments, left_index=True, right_index=True)
      name  department_id  department
id
1  Dilbert           1  Engineering
2   Wally           1           HR
3  Catbert           2  Management

```

Merging on the index is the default behaviour of the `join()` method:

```
>>> employees.join(departments)
      name  department_id  department
id
1  Dilbert              1  Engineering
2    Wally              1           HR
3  Catbert              2   Management
```

Merging on indices and columns can also be mixed, specifying either the `left_on/right_index` or the `left_index/right_on` parameter pairs:

```
>>> employees = pd.DataFrame({
    'id': [1, 2, 3],
    'name': ['Dilbert', 'Wally', 'Catbert'],
    'department_id': [1, 1, 2]})
>>> employees
   id  name  department_id
0   1  Dilbert            1
1   2   Wally            1
2   3  Catbert            2

>>> departments = pd.DataFrame({
    'id': [1, 2, 3],
    'department': ['Engineering', 'HR', 'Management']})
>>> departments = departments.set_index('id')
>>> departments
      department
id
1  Engineering
2           HR
3   Management

>>> pd.merge(employees, departments, left_on='department_id', right_index=True)
   id  name  department_id  department
0   1  Dilbert            1  Engineering
1   2   Wally            1  Engineering
2   3  Catbert            2           HR
```

The type of the join to be performed in terms of set arithmetic can be defined using the `how` keyword. The default option is `inner`; only entries common to both input datasets are contained in the result:

```
>>> employees = pd.DataFrame({
    'employee': ['Dilbert', 'Pointy Haired Boss', 'Dogbert'],
    'department': ['Engineering', 'Management', 'Evil Operations']})
>>> employees
   employee      department
0    Dilbert  Engineering
1 Pointy Haired Boss  Management
2    Dogbert  Evil Operations

>>> departments = pd.DataFrame({
    'department': ['Engineering', 'Management', 'Marketing'],
    'location': ['basement', 'upper floor', 'middle floor']})
>>> departments
   department      location
0  Engineering  basement
1  Management  upper floor
2  Marketing    middle floor
```

```
0 Engineering      basement
1 Management      upper floor
2 Marketing        middle floor
```

```
>>> pd.merge(employees, departments, how='inner')
      employee  department  location
0      Dilbert  Engineering  basement
1 Pointy Haired Boss  Management  upper floor
```

The option outer fills up missing entries (i.e. entries not common to both input datasets) with NaN in the result:

```
>>> pd.merge(employees, departments, how='outer')
      employee  department  location
0      Dilbert  Engineering  basement
1 Pointy Haired Boss  Management  upper floor
2      Dogbert  Evil Operations  NaN
3      NaN      Marketing  middle floor
```

The options left and right preserve all values from the left resp. right side, and fill up all the missing entries on the other side with NaN:

```
>>> pd.merge(employees, departments, how='left')
      employee  department  location
0      Dilbert  Engineering  basement
1 Pointy Haired Boss  Management  upper floor
2      Dogbert  Evil Operations  NaN
```

```
>>> pd.merge(employees, departments, how='right')
      employee  department  location
0      Dilbert  Engineering  basement
1 Pointy Haired Boss  Management  upper floor
2      NaN      Marketing  middle floor
```

If the two input datasets have columns with the same name that are not used to perform the join operation, a suffix (_x and _y) is added to both columns to prevent conflicts:

```
>>> employees.index.names = ['id']
>>> employees = employees.reset_index()
>>> employees
   id  employee  department
0  0      Dilbert  Engineering
1  1 Pointy Haired Boss  Management
2  2      Dogbert  Evil Operations

>>> departments.index.names = ['id']
>>> departments = departments.reset_index()
>>> departments
   id  department  location
0  0  Engineering  basement
1  1  Management  upper floor
2  2  Marketing  middle floor

>>> pd.merge(employees, departments, on='department')
   id_x  employee  department  id_y  location
0     0      Dilbert  Engineering     0  basement
1     1 Pointy Haired Boss  Management     1  upper floor
```

A list of custom suffixes can be set using the `suffixes` parameter:

```
>>> pd.merge(employees, departments, on='department', suffixes=['_emp', '_dep'])
   id_emp  employee  department  id_dep  location
0        0      Dilbert  Engineering        0  basement
1        1  Pointy Haired Boss  Management        1  upper floor
```

3.8 Aggregation

Computing aggregations is an essential technique for efficient summarization of data sets. The planets dataset of the seaborn package is useful for practicing aggregations:

```
>>> import seaborn as sns
>>> planets = sns.load_dataset('planets')
```

A good starting point is to get an overview over the dataset using the `describe()` function, which is a convenience method that performs a couple of aggregations for the purpose of understanding rather than further processing the data:

```
>>> planets.describe()
   number  orbital_period  mass  distance  year
count  1035.000000      992.000000  513.000000  808.000000  1035.000000
mean    1.785507      2002.917596   2.638161  264.069282  2009.070531
std     1.240976     26014.728304   3.818617  733.116493   3.972567
min     1.000000       0.090706   0.003600   1.350000  1989.000000
25%     1.000000       5.442540   0.229000   32.560000  2007.000000
50%     1.000000      39.979500   1.260000   55.250000  2010.000000
75%     2.000000     526.005000   3.040000  178.500000  2012.000000
max     7.000000    730000.000000  25.000000  8500.000000  2014.000000
```

Important aggregation functions are:

Function	Returns
<code>count()</code>	number of entries (NaN not counted)
<code>min()</code>	minimum value
<code>max()</code>	maximum value
<code>sum()</code>	sum (addition)
<code>prod()</code>	product (multiplication)
<code>mean()</code>	mean (arithmetic average)
<code>median()</code>	median (middle value)
<code>std()</code>	standard deviation
<code>var()</code>	variance
<code>mad()</code>	mean absolute deviation

Aggregations on a DataFrame result in summarized columns. To aggregate rows instead of columns, the `axis` parameter can be set accordingly:

```
>>> planets.mean(axis='columns')
```

The `axis` parameters describe what is to be aggregated (the *columns* of each row), not what the result should be!

3.9 Grouping

Grouping allows to split a dataset up based on its values or index, perform computations within the groups and combine the group results together to overall results. Grouping is a three-step process:

1. split: breaking up and grouping a DataFrame (based on the values of a specified key or other property)
2. apply: perform computations within each group:
 1. filter: remove or retain values for further processing
 2. transform: map the input values to output values
 3. aggregate: reduce the multitude of values to a single value (or a smaller amount of values)
 4. apply: perform computations on the aggregation result(s)
3. combine: merge the results to a single resulting dataset

The `groupby()` method allows to perform those three steps together in an efficient way. When called on a DataFrame, it returns a `DataFrameGroupBy` object, which is a special (grouped) view onto the underlying DataFrame:

```
>>> import seaborn as sns
>>> planets = sns.load_dataset('planets')
>>> planets.groupby('year')
<pandas.core.groupby.groupby.DataFrameGroupBy object at 0x7f9db32f2eb8>
```

A `DataFrameGroupBy` is a collection of DataFrames that allows for the operations filter, transform, aggregate and apply. No computation is performed until an aggregation is applied (lazy evaluation), which returns a new DataFrame:

```
>>> planets.groupby('year').sum()
   number  orbital_period      mass  distance
year
1989      1      83.888000  11.68000    40.57
1992      6      91.803900   0.00000     0.00
1994      3      98.211400   0.00000     0.00
...
```

Selecting a column on a `DataFrameGroupBy` object returns a `SeriesGroupBy` object, which can be also used for aggregations and the like:

```
>>> planets.groupby('year')['distance']
<pandas.core.groupby.groupby.SeriesGroupBy object at 0x7f9db3224f60>
```

A `GroupBy` object allows to iterate over the individual groups, yielding the group key and the DataFrame:

```
>>> for (key, df) in planets.groupby('year'):
    print(key, ', '.join(df.columns))
1989 method, number, orbital_period, mass, distance, year
1992 method, number, orbital_period, mass, distance, year
1994 method, number, orbital_period, mass, distance, year
...
```

However, the `apply()` method is usually faster and more convenient than an explicit iteration.

When a method of a DataFrame is called on a `GroupBy` object, it is dispatched to each of the underlying DataFrame objects:

```
>>> planets.groupby('year').first()
              method  number  orbital_period    mass  distance
year
1989  Radial Velocity      1      83.888000  11.6800    40.57
1992   Pulsar Timing      3      25.262000     NaN     NaN
1994   Pulsar Timing      3      98.211400     NaN     NaN
...
```

As mentioned earlier, after grouping and before combining the data, different operations can be performed on the grouped data.

The `filter()` method executes a predicate function (or lambda expression) on every entry, retains it in the dataset (matching condition) or discards it from the dataset (not matching condition). The predicate function/lambda expression expects a DataFrame and returns a boolean:

```
>>> teams = ['Mercedes', 'Mercedes', 'Ferrari', 'Ferrari']
>>> drivers = ['Hamilton', 'Bottas', 'Vettel', 'Raikkoennen']
>>> points = [408, 247, 320, 251]
>>> championship = df.DataFrame(
    {'team': teams, 'driver': drivers, 'points': points})
>>> championship
   team      driver  points
0  Mercedes  Hamilton    408
1  Mercedes   Bottas    247
2   Ferrari   Vettel    320
3   Ferrari Raikkoennen  251

>>> championship.groupby('team').filter(lambda x: x['points'].mean() > 300)
   team      driver  points
0  Mercedes  Hamilton    408
1  Mercedes   Bottas    247
```

The DataFrame is grouped by team. For every team the mean of points scored is calculated, and only entries with a team's point mean above 300 are retained. This filtering uses a predicate function:

```
>>> def below_600(x):
    return x['points'].sum() < 600
>>> championship.groupby('team').filter(below_600)
   team      driver  points
2  Ferrari   Vettel    320
3  Ferrari Raikkoennen  251
```

The `transform()` method allows to map the input data record by record to output data of the same shape:

```
>>> championship.groupby('team')['points'].transform(lambda x: x / x.mean())
0    1.245802
1    0.754198
2    1.120841
3    0.879159
```

Each driver's ratio of points scored to the team is computed in terms of mean points per team. Notice that the points column was selected, so `x` refers to a Series, not to a DataFrame.

The `aggregate()` method allows to reduce a group in two fundamental ways:

First, by applying one or more aggregation functions that are passed either as a function or as a function name (string):

```
>>> championship.groupby('team').aggregate([min, 'max'])
              driver      points
              min      max    min  max
team
Ferrari    Raikkoennen  Vettel    251  320
Mercedes      Bottas  Hamilton    247  408
```

Second, by applying different aggregation functions for each column, by providing a dictionary that maps a function to every column:

```
>>> championship['position'] = [1, 5, 2, 3]
>>> championship.groupby('team').aggregate({'points': max, 'position': min})
              points  position
team
Ferrari          320          2
Mercedes          408          1
```

The `apply()` method allows to execute a function on every group result. It takes a `DataFrame`/`Series` and returns either a `DataFrame`/`Series` object, or the function reduces the group results further to a single scalar:

```
>>> championship.groupby('team')['points'].apply(sum)
team
Ferrari    571
Mercedes   655
Name: points, dtype: int64
```

The grouping of the data is not limited to a single column name. Different alternatives are available.

First, provide a list/array/series/index of group keys, telling every entry in which group to go:

```
>>> names = ['Harry Potter', 'Draco Malfoy', 'Hermine Granger', 'Ron Weasley']
>>> students = pd.Series(names)
>>> houses = ['Griffindor', 'Slytherin', 'Griffindor', 'Griffindor']
>>> students.groupby(houses).apply(lambda s: ', '.join(s))
Griffindor    Harry Potter, Hermine Granger, Ron Weasley
Slytherin                                Draco Malfoy
dtype: object
```

Second, provide a dictionary that maps the index keys to groups:

```
>>> courses = ['Math', 'English', 'History', 'Geography', 'Music', 'Biology']
>>> results = ['A', 'C', 'E', 'B', 'D', 'F']
>>> grouping = {'A': 'good', 'B': 'good', 'C': 'ok', 'D': 'ok', 'E': 'bad', 'F': 'bad'}
>>> marks = pd.DataFrame({'course': courses, 'result': results})
>>> marks = marks.set_index('result')
>>> marks
              course
result
A              Math
C              English
E              History
B          Geography
D              Music
```


F Biology

```
>>> marks.groupby(grouping).aggregate(lambda c: ', '.join(c))
               course
bad   History, Biology
good   Math, Geography
ok     English, Music
```

Third, provide any function that maps a input (index) to a output (group):

```
>>> lectures = ['Math: Calculus', 'Math: Statistics',
                'Computer Science: Algorithms', 'Computer Science: Data Structures']
>>> professors = ['Smith', 'Myers', 'Dijkstra', 'Kernighan']
>>> plan = pd.DataFrame({'lecture': lectures, 'professor': professors})
>>> plan = plan.set_index('lecture')
>>> plan
```

	professor
lecture	
Math: Calculus	Smith
Math: Statistics	Myers
Computer Science: Algorithms	Dijkstra
Computer Science: Data Structures	Kernighan

```
>>> plan.groupby(lambda l: l.split(':')[0]).aggregate(lambda p: ', '.join(p))
               professor
Computer Science  Dijkstra, Kernighan
Math              Smith, Myers
```

And fourth, use a combination thereof, which results in a MultiIndex:

```
>>> marks.groupby([str.lower, grouping]).aggregate(lambda m: ' '.join(m))
               course
a good        Math
b good  Geography
c ok          English
d ok          Music
e bad    History
f bad    Biology
```

3.10 Pivot Tables

Pivot Tables are essentially a multidimensional version of the GroupBy aggregation. A DataFrame can be analyzed in two dimensions. In terms of GroupBy, the split and combine steps are performed along a two-dimensional grid, and the two dimensions can be defined (as index and columns).

The “titanic” dataset of the Seaborn package is a good example for a multidimensional analysis. This GroupBy operation aggregates the survival rates by both sex *and* class:

```
>>> titanic.groupby(['sex', 'class'])['survived'].aggregate('mean').unstack()
class      First      Second      Third
sex
female  0.968085  0.921053  0.500000
male    0.368852  0.157407  0.135447
```

The instruction reads as “group by sex and class, select the survived column, calculate the mean thereof, and display the result in a two-dimensional view”.

The same result can be achieved with less typing using the `pivot_table()` method:

```
>>> titanic.pivot_table('survived', index='sex', columns='class')
class      First      Second      Third
sex
female  0.968085  0.921053  0.500000
male    0.368852  0.157407  0.135447
```

Calculating the mean is the default aggregation of the `pivot_table()` method. The instruction reads as “calculate the mean of the survived column by sex and class”.

Grouping is not restricted to single values. More dimensions can be brought in by providing a list of criteria.

The `cut()` method categorizes a series of values using the given boundaries. The age categories are then used as an additional (third) dimension:

```
>>> age = pd.cut(titanic['age'], [0, 18, 80])
>>> titanic.pivot_table('survived', ['sex', age], 'class')
class      First      Second      Third
sex  age
female (0, 18]  0.909091  1.000000  0.511628
      (18, 80]  0.972973  0.900000  0.423729
male   (0, 18]  0.800000  0.600000  0.215686
      (18, 80]  0.375000  0.071429  0.133663
```

The `qcut()` method splits up a series of values to the given number of quantiles. The fare quantiles are then used as an additional (fourth) dimension:

```
>>> fare = pd.qcut(titanic['fare'], 2)
>>> titanic.pivot_table('survived', ['sex', age], [fare, 'class'])
fare      (-0.001, 14.454]      (14.454, 512.329]
class      First      Second      Third      First      Second      Third
sex  age
female (0, 18]      NaN  1.000000  0.714286      0.909091  1.000000  0.318182
      (18, 80]      NaN  0.880000  0.444444      0.972973  0.914286  0.391304
male   (0, 18]      NaN  0.000000  0.260870      0.800000  0.818182  0.178571
      (18, 80]      0.0  0.098039  0.125000      0.391304  0.030303  0.192308
```

The `pivot_table()` method has a lot of additional parameters. Its signature looks as follows:

```
DataFrame.pivot_table(values=None, index=None, columns=None, aggfunc='mean',
                        fill_values=None, margins=False, dropna=True,
                        margins_name='All')
```

The parameters have the following meaning:

- `values`: the column of interest (to be aggregated)
- `index`: the y-axis group keys
- `columns`: the x-axis group keys
- `aggfunc`: the aggregation to be performed on values
 - accepts either a list of functions
 - or a dictionary specifying column/aggregation pairs (values can be omitted)
- `fill_value`: value to use for empty fields
- `margins`: whether or not to compute totals

- `dropna`: whether or not to ignore NaN entries
- `margins_name`: labels for the margin totals (default: 'All')

Example:

```
>>> titanic.pivot_table(values='survived', index='embark_town', columns='alone',
                        aggfunc='mean', fill_value=False, margins=True,
                        dropna=True, margins_name='survival rate')
alone                False      True  survival rate
embark_town
Cherbourg      0.674699  0.435294      0.553571
Queenstown    0.350000  0.403509      0.389610
Southampton   0.462151  0.256997      0.336957
survival rate  0.505650  0.300935      0.382452
```

3.11 Vectorized String Operations

Real-world datasets often contain a lot of messy string data. Pandas supports vectorized string operations that can easily be applied on entire columns or datasets without worrying about the shape of the data or missing values. Vectorized operations are also more efficient than explicitly iterating over the values and calling the operation on each value.

Series and Index objects have a `str` attribute that provides functionality to deal with the underlying strings. (A column of a DataFrame is a Series and therefore also has a `str` attribute.)

Pandas implements a good deal of Python's native string and regular expression functions as methods of the `str` attribute, which are demonstrated on the following dataset:

```
>>> names = ['Dilbert', 'Alice', 'Wally', 'Pointy Haired Boss']
>>> notes = ['nerdy, whiny', 'aggressive, grumpy', 'lazy, dorky', 'clueless, cocky']
>>> review = pd.DataFrame({'employees': names, 'properties': notes})
>>> review
   employees      properties
0    Dilbert  nerdy, whiny
1     Alice  aggressive, grumpy
2     Wally    lazy, dorky
3 Pointy Haired Boss  clueless, cocky
```

Predicate methods check a property of a string and return a boolean value indicating whether or not the property in question applies to it:

Method	Description
<code>startswith(prefix)</code>	begins with prefix?
<code>endswith(suffix)</code>	begins with suffix?
<code>isalnum()</code>	consists of letters and digits only?
<code>isalpha()</code>	consists of letters only?
<code>isdigit()</code>	consists of digits only? (like 3, 2 ²)
<code>isnumeric()</code>	is a numeric expression? (like ½, 2 ²)
<code>isdecimal()</code>	is a numeric expression? (like 123)
<code>isspace()</code>	consists of spaces only?
<code>istitle()</code>	is every word written in title case?
<code>islower()</code>	consists of lower case letters only?
<code>isupper()</code>	consists of upper case letters only?

These methods perform a transformation on the underlying string and return the result of that transformation:

Method	Description
<code>ljust(width)</code>	left align to width
<code>rjust(width)</code>	right align to width
<code>center(width)</code>	center align to width
<code>pad(width, side)</code>	justify to width with side ('left', 'right', 'both')
<code>zfill(width)</code>	fill up with 0 from left to width
<code>strip()</code>	remove trailing whitespace
<code>rstrip()</code>	remove trailing whitespace on the left
<code>rstrip()</code>	remove trailing whitespace on the right
<code>wrap(n)</code>	add newline after n characters
<code>join(s)</code>	separate characters with string s
<code>cat()</code>	concatenate the strings
<code>upper()</code>	all upper case letters
<code>lower()</code>	all lower case letters
<code>capitalize()</code>	first letter of first word upper case
<code>swapcase()</code>	upper to lower, and lower to upper case
<code>translate(table)</code>	apply map of translation rules in table
<code>normalize(form)</code>	'NFC', 'NFKC', 'NFD' or 'NFKD' unicode normalization
<code>repeat(n)</code>	repeats the string n times
<code>slice_replace(a, z, repl)</code>	replaces the slice [a:z] with repl
<code>get(i)/[i]</code>	get character at index i
<code>slice(a, z, s)/[a:z:s]</code>	slice (from a to z with step s)

The `translate` method requires a table, which can be created using the string method `maketrans`:

```
>>> table = str.maketrans({'t': 'th', 'i': 'y'})
>>> review['employees'].str.translate(table)
0          Dylberth
1          Alyce
2          Wally
3    Poynthly Hayred Boss
```

The following miscellaneous methods return neither a boolean value nor a modified string, but either a number or other data structure:

Method	Description
<code>len()</code>	length in characters
<code>find(s)</code>	start index of substring s (-1 if not contained)
<code>rfind(s)</code>	like <code>find()</code> , but starts from the end
<code>index(s, a, z)</code>	like <code>find()</code> with range a:z (ValueError if not contained)
<code>rindex(s, a, z)</code>	like <code>index()</code> , but starts from the end
<code>partition(sep)</code>	split into three parts: before, sep, after (default sep: whitespace)
<code>rpartition(sep)</code>	like <code>partition()</code> , but starts from the end
<code>get_dummies(sep)</code>	transform encoded string into DataFrame using sep to split values

The `get_dummies()` method is especially useful when meaning is encoded into a string using multiple, separated values:

```
>>> review['properties'].str.get_dummies(',')
```

	aggressive	clueless	cocky	dorky	grumpy	lazy	nerdy	whiny
0	0	0	0	0	0	0	1	1
1	1	0	0	0	1	0	0	0
2	0	0	0	1	0	1	0	0
3	0	1	1	0	0	0	0	0

These methods implement functionality from Python's regular expression library (re):

Method	Description
<code>match(pat)</code>	does the pattern <code>pat</code> match? (see <code>re.match</code>)
<code>contains(str)</code>	is the string <code>str</code> contained? (see <code>re.search</code>)
<code>extract(pat)</code>	extracts the groups from the pattern <code>pat</code>
<code>findall(pat)</code>	returns all occurrences matching <code>pat</code>
<code>replace(pat, repl)</code>	replaces occurrences of <code>pat</code> with <code>repl</code>
<code>count(pat)</code>	number of matches of <code>pat</code>
<code>split(pat)</code>	split at matches of <code>pat</code>
<code>rsplit(pat)</code>	like <code>split()</code> , but starts from the end

3.12 Time Series

Pandas has strong capabilities to deal with dates, times and data indexed by date and time. The notion of time can be expressed in different concepts:

- *Time stamps* refer to a particular moment, like June 24th 1987, 8:25 a.m.
- *Time intervals* and *periods* express a length of time between a beginning and an end point, like the year 2019 or the second week of 2019.
 - *Periods* are a special kind of interval: They do not overlap with other intervals and are of uniform length, like a day or an hour.
- *Time deltas* or *durations* express an exact length of time, like 9.87 seconds.

Pandas capabilities for dealing with date and time set up on Python's native date and time tools.

Python's built-in datetime module with the datetime type is useful for expressing single dates:

```
>>> from datetime import datetime
>>> birth = datetime(year=1987, month=6, day=24, hour=8, minute=25)
>>> birth
datetime.datetime(1987, 6, 24, 8, 25)

>>> birth.strftime('%A') # %A: day of week
'Wednesday'
```

The third-party dateutil module can parse dates of various string formats:

```
>>> from dateutil import parser
>>> birth = parser.parse("24th of June, 1987 at 8:25 a.m")
>>> birth
datetime.datetime(1987, 6, 24, 8, 25)

>>> birth.strftime('%A') # %A: day of week
'Wednesday'
```

The third-party `pytz` module helps to deal with time zones.

Those tools are convenient, but do not scale for big data sets consisting of date and time information. One alternative is NumPy's `datetime64` type.

A better alternative in the context of Pandas is the `Timestamp` object, which combines the comfort of Python's native `datetime` and third-party `dateutil` with the efficiency of NumPy's `datetime64`.

Dates can be parsed as with `dateutil`:

```
>>> birth = pd.to_datetime("24th of June, 1987 at 8:25 a.m.")
>>> birth
Timestamp('1987-06-24 08:25:00')
```

```
>>> birth.strftime('%A')
'Wednesday'
```

Vectorized operations on dates can be performed as efficiently as with NumPy's `datetime64` type:

```
>>> date = pd.to_datetime("1st of January 2019")
>>> date + pd.to_timedelta(range(3), 'D')
DatetimeIndex(['2019-01-01', '2019-01-02', '2019-01-03'], dtype='datetime64[ns]', freq=None)
```

A `DatetimeIndex` is used to index `Timestamp` objects in a `Series` or `DataFrame`. It offers powerful slicing and indexing operations:

```
>>> index = pd.DatetimeIndex(['2015-01-01', '2016-04-01', '2017-07-01', '2018-10-01'])
>>> dates = pd.Series(range(4), index=index)
>>> dates
2015-01-01    0
2016-04-01    1
2017-07-01    2
2018-10-01    3
dtype: int64
```

```
>>> dates['2016-01-01': '2017-12-31'] # slicing
2016-04-01    1
2017-07-01    2
dtype: int64
```

```
>>> dates['2016'] # indexing
2016-04-01    1
dtype: int64
```

Pandas implements the different time concepts with different data types and indices:

Concept	Type	Index Type	Python/NumPy Type
Time Stamp	Timestamp	DatetimeIndex	datetime/datetime64
Time Period	Period	PeriodIndex	-/datetime64
Time Delta/Duration	Timedelta	TimedeltaIndex	timedelta/timedelta64

These types and indices can be used directly, but Pandas offers convenience functions for easier parsing and handling of entire `Series`.

The `pd.to_datetime()` function yields a `Timestamp` if a single date is passed, and a `DatetimeIndex` if a series of dates (in any format) is passed:

```
>>> date pd.to_datetime('2018-12-24')
>>> date
Timestamp('2018-12-24 00:00:00')

>>> index = pd.to_datetime(['2018-03-17', '25th of March 1992',
                             datetime(2019, 6, 24), '1984-Jul-20', '20190101'])
DatetimeIndex(['2018-03-17', '1992-03-25', '2019-06-24', '1984-07-20',
               '2019-01-01'],
              dtype='datetime64[ns]', freq=None)
```

A DatetimeIndex can be converted to a PeriodIndex using the `to_period()` method by indicating a frequency code, like 'D' for days:

```
>>> periods = index.to_period('D')
PeriodIndex(['2018-03-17', '1992-03-25', '2019-06-24', '1984-07-20',
            '2019-01-01'],
           dtype='period[D]', freq='D')
```

A timedeltaIndex, describing the difference between dates, can be created by a subtraction, for example:

```
>>> deltas = index - index[0]
>>> deltas
TimedeltaIndex(['0 days', '-9488 days', '464 days', '-12293 days', '290 days'],
               dtype='timedelta64[ns]', freq=None)
```

3.12.1 Sequences

Pandas offers convenience functions to create regular date sequences. Like Python's `range()` and NumPy's `np.arange()`, they accept a beginning and end point, and an optional frequency.

A sequence of dates can be created using the `pd.date_range()` function:

```
>>> pd.date_range('2018-01-01', '2018-01-08')
DatetimeIndex(['2018-01-01', '2018-01-02', '2018-01-03', '2018-01-04',
               '2018-01-05', '2018-01-06', '2018-01-07', '2018-01-08'],
              dtype='datetime64[ns]', freq='D')
```

Instead of defining an end date, the number of periods can be defined:

```
>>> pd.date_range('2018-01-01', periods=8)
DatetimeIndex(['2018-01-01', '2018-01-02', '2018-01-03', '2018-01-04',
               '2018-01-05', '2018-01-06', '2018-01-07', '2018-01-08'],
              dtype='datetime64[ns]', freq='D')
```

Any combination of two indications (start, end, frequency) is enough to create a sequence:

```
>>> pd.date_range(start='2018-01-01', end='2018-01-08') # start and end
DatetimeIndex(['2018-01-01', '2018-01-02', '2018-01-03', '2018-01-04',
               '2018-01-05', '2018-01-06', '2018-01-07', '2018-01-08'],
              dtype='datetime64[ns]', freq='D')

>>> pd.date_range(start='2018-01-01', periods=8) # start and periods
DatetimeIndex(['2018-01-01', '2018-01-02', '2018-01-03', '2018-01-04',
               '2018-01-05', '2018-01-06', '2018-01-07', '2018-01-08'],
              dtype='datetime64[ns]', freq='D')
```

```
>>> pd.date_range(end='2018-01-08', periods=8) # end and periods
DatetimeIndex(['2018-01-01', '2018-01-02', '2018-01-03', '2018-01-04',
               '2018-01-05', '2018-01-06', '2018-01-07', '2018-01-08'],
              dtype='datetime64[ns]', freq='D')

>>> pd.date_range(start='2018-01-01', end='2018-01-08', periods=4) # all three
DatetimeIndex(['2018-01-01 00:00:00', '2018-01-03 08:00:00',
               '2018-01-05 16:00:00', '2018-01-08 00:00:00'],
              dtype='datetime64[ns]', freq=None)
```

The frequency defaults to one day. In the last example, where start, end *and* periods were given, no fixed frequency is used, but calculated to evenly distribute the dates between start and end.

A frequency can be defined using the `freq` parameter:

```
>>> pd.date_range(start='2018-01-01', periods=4, freq='H')
DatetimeIndex(['2018-01-01 00:00:00', '2018-01-01 01:00:00',
               '2018-01-01 02:00:00', '2018-01-01 03:00:00'],
              dtype='datetime64[ns]', freq='H')

>>> pd.date_range(start='2018-01-01', periods=4, freq='M')
DatetimeIndex(['2018-01-31', '2018-02-28', '2018-03-31', '2018-04-30',
               '2018-05-31', '2018-06-30', '2018-07-31', '2018-08-31'],
              dtype='datetime64[ns]', freq='M')
```

Regular sequences of periods can be created using the `period_range()` function:

```
>>> pd.period_range('2018-01', periods=12, freq='M')
PeriodIndex(['2018-01', '2018-02', '2018-03', '2018-04', '2018-05', '2018-06',
             '2018-07', '2018-08', '2018-09', '2018-10', '2018-11', '2018-12'],
            dtype='period[M]', freq='M')
```

Regular sequences of durations/time deltas can be created using the `timedelta_range()` function:

```
>>> pd.timedelta_range(0, periods=10, freq='H')
TimedeltaIndex(['00:00:00', '01:00:00', '02:00:00', '03:00:00', '04:00:00',
                '05:00:00', '06:00:00', '07:00:00', '08:00:00', '09:00:00'],
               dtype='timedelta64[ns]', freq='H')
```

Pandas offers the following *date* frequencies (at either the start or end of each period):

Code	Frequency	Code	Frequency
AS	year start	A	year end
BAS	business year start	BA	business year end
QS	quarter start	Q	quarter end
BQS	business quarter start	BQ	business quarter end
MS	month start	M	month end
BMS	business month start	BM	business month end

And these *time* frequencies:

Code	Frequency	Code	Frequency
W	week	T	minute

Code	Frequency	Code	Frequency
D	day	S	second
B	business day	L	millisecond
H	hour	U	microsecond
BH	business hour	N	nanosecond

Quarter and year frequencies can be marked with a month suffix, weekly frequencies can be marked with a day suffix in order to specify the split points:

```
>>> pd.date_range('2018-01-01', periods=8, freq='QS-JAN')
DatetimeIndex(['2018-01-01', '2018-04-01', '2018-07-01', '2018-10-01',
               '2019-01-01', '2019-04-01', '2019-07-01', '2019-10-01'],
              dtype='datetime64[ns]', freq='QS-JAN')

>>> pd.date_range('2018-01-01', periods=8, freq='AS-JUL')
DatetimeIndex(['2018-07-01', '2019-07-01', '2020-07-01', '2021-07-01',
               '2022-07-01', '2023-07-01', '2024-07-01', '2025-07-01'],
              dtype='datetime64[ns]', freq='AS-JUL')

>>> pd.date_range('2018-01-01', periods=8, freq='W-SUN')
DatetimeIndex(['2018-01-07', '2018-01-14', '2018-01-21', '2018-01-28',
               '2018-02-04', '2018-02-11', '2018-02-18', '2018-02-25'],
              dtype='datetime64[ns]', freq='W-SUN')
```

The frequency codes refer to instances of the module `pandas.tseries.offsets` and can be used as functions:

```
>>> pd.date_range('2018-01-01', periods=8, freq=BDay())
DatetimeIndex(['2018-01-01', '2018-01-02', '2018-01-03', '2018-01-04',
               '2018-01-05', '2018-01-08', '2018-01-09', '2018-01-10'],
              dtype='datetime64[ns]', freq='B')
```

Frequency codes can be combined with additional numbers to create custom periods, such as 1 hour and 45 minutes:

```
>>> pd.date_range('2018-01-01', periods=8, freq='23H15T')
DatetimeIndex(['2018-01-01 00:00:00', '2018-01-01 23:15:00',
               '2018-01-02 22:30:00', '2018-01-03 21:45:00',
               '2018-01-04 21:00:00', '2018-01-05 20:15:00',
               '2018-01-06 19:30:00', '2018-01-07 18:45:00'],
              dtype='datetime64[ns]', freq='1395T')
```

3.12.2 Resampling, Shifting, Windowing

Resampling, Shifting and Windowing are useful operations to analyze time series. Analyzing stock prices is a important use case, and stock prices can be conveniently loaded with the `pandas-datareader` package from Yahoo Finance, for example the closing price of the Microsoft stock:

```
>>> from pandas_datareader import data
>>> msft = data.DataReader('MSFT', start='1986', end='2019', data_source='yahoo')
>>> msft = msft['Close']
>>> msft.describe()
count      8269.000000
```

```

mean      25.047959
std       22.397970
min        0.090278
25%        2.992188
50%       25.930000
75%       32.345001
max       115.610001
Name: Close, dtype: float64

```

The stock price over time can be visualized using the matplotlib library, using the opticts from the seaborn package:

```

>>> import matplotlib.pyplot as plt
>>> import seaborn
>>> seaborn.set()
>>> msft.plot();
>>> plt.show();

```

The time series can be resampled to a higher or lower frequency using the `resample()` method, which can be used to perform a data aggregation. The simpler `asfreq()` converts the frequency by simply selecting data (as opposed to aggregating them).

Both methods are used here to visualize the stock price by business year compared to the daily closing prices:

```

>>> msft.plot(style='-', alpha=0.5)
>>> msft.resample('BA').mean().plot(style=':') # mean of business year
>>> msft.asfreq('BA').plot(style='--') # business year's closing price
>>> plt.legend(['original', 'resample', 'asfreq'], loc='upper left')
>>> plt.show()

```

Time shifts are useful to compute differences over time. The method `tshift()` can be used to shift the index values, whereas the method `shift()` shifts the data itself. The shift is specified in multiples of the underlying frequency:

```

>>> cs = data.DataReader('CS', start='2000', end='2019', data_source='yahoo')
>>> cs = cs['Close'].asfreq('D')
>>> cs.plot()
>>> cs.shift(365).plot()
>>> plt.legend(['original', 'shift(365)'], loc='upper left')
>>> plt.show()

```

Rolling statistics can be used to perform different aggregations over a rolling data window, like the mean of the last 365 days relative to every day.

```

>>> aapl = data.DataReader('AAPL', start='2000', end='2019', data_source='yahoo')
>>> aapl = aapl['Close']
>>> rolling = aapl.rolling(365, center=True)
>>> aapl.plot()
>>> rolling.mean().plot()
>>> plt.legend(['original', 'mean over 365 days'], loc='upper left')
>>> plt.show()

```




Figure 1: Microsoft Stock Price



Figure 2: Resampling and Frequency Conversion




Figure 3: Shifting



Figure 4: Rolling Window

3.13 High-Performance Pandas: eval() and query()

Even though vectorized operations in NumPy and Pandas are much more efficient than explicit iterations, compound expressions still cause a big memory overhead to store the intermediate steps.

Consider this masking operation:

```
>>> mask = (x > 0.5) & (y < 0.5)
```

Every intermediate step allocates memory, which becomes more obvious if the above expression is written as such:

```
>>> tmp1 = (x > 0.5)
>>> tmp2 = (y < 0.5)
>>> mask = tmp1 & tmp2
```

Pandas eval() and query() methods, which are based on the [Numexpr](#) package, can do without full-sized temporary arrays and hence are much lighter on memory consumption than vectorized operations.

The eval() function accepts a string expression describing an operation on DataFrames:

```
>>> %load_ext memory_profiler
>>> n = 100_000_000
>>> cols = 10
>>> df1, df2, df3, df4 = (pd.DataFrame(np.random.random(n).reshape(n//cols, cols))
                        for i in range(4))
>>> %memit df1 + df2 + df3 + df4 # vectorized operation
peak memory: 5175.80 MiB, increment: 1450.38 MiB

>>> %memit sum = pd.eval('df1 + df2 + df3 + df4') # numeric expression
peak memory: 4323.87 MiB, increment: 1159.55 MiB
```

Supported are arithmetic (+, -, *, /), comparison (==, !=, >, >=, <, <=), bitwise resp. element-wise (&, |) and logical (and, or) operators, as well as indexing (df['col']) and attribute access (df.attr). Constructs like loops and function calls aren't available with eval(), but need direct use of the Numexpr package.

DataFrame as its own eval() method. In addition to the features of the pd.eval() function, it supports direct column access by their names and access to variables:

```
>>> n = 3_000_000
>>> df = pd.DataFrame(np.random.random(n).reshape(n//3, 3),
                      columns=['A', 'B', 'C'])

>>> %memit (df['A'] + df['B']) / (df['C'] - 1) # vectorized operation
peak memory: 123.24 MiB, increment: 19.62 MiB

>>> %memit pd.eval('(df.A + df.B) / (df.C - 1)') # columns as attributes
peak memory: 112.56 MiB, increment: 8.15 MiB

>>> %memit df.eval('(A + B) / (C - 1)') # direct column access
peak memory: 143.60 MiB, increment: 8.21 MiB

>>> %memit df.eval('D = (A + B) / (C - 1)', inplace=True) # create new column
peak memory: 166.62 MiB, increment: 30.73 MiB
```

```
>>> %memit df.eval('D = (A + B) / C', inplace=True)
peak memory: 166.66 MiB, increment: 0.00 MiB # overwrite existing column
```

Variables from the enclosing scope can be used with the @ prefix (in order to distinguish them from columns):

```
>>> mean = df['A'].mean()
>>> %memit df.eval('D = (A + B) / (C - @mean)')
peak memory: 227.79 MiB, increment: 61.04 MiB
```

Masking and filtering expressions cannot be expressed using the DataFrame.eval() method. The method DataFrame.query() makes this possible:

```
>>> %memit df[(df.A > mean) & (df.B < mean)] # vectorized operation
peak memory: 228.57 MiB, increment: 0.00 MiB
```

```
>>> %memit pd.eval('df[(df.A > mean) & (df.B < mean)]') # pd.eval()
peak memory: 230.36 MiB, increment: 1.88 MiB
```

```
>>> %memit df.query('A > @mean and B < @mean') # DataFrame.query()
peak memory: 230.98 MiB, increment: 0.00 MiB
```

Notice that the bitwise (element-wise) & operator has to be translated to and in the expression for the query() method.

eval() and query() have some downsides:

1. They deal with strings as opposed to Python syntax, which makes it harder to detect syntax errors for both the human eye and tools.
2. They have some computational overhead, which might outweigh the possible savings on temporary memory usage by far.

A good starting point in the decision between vectorized operations and eval()/query() is the size of a DataFrame:

```
>>> n = 300_000
>>> df = pd.DataFrame(np.random.random(n).reshape(n//3, 3),
                      columns=['A', 'B', 'C'])
>>> df.values.nbytes / (1024*1024) # size in megabytes
2.288818359375
```

If a DataFrame doesn't fit into the CPU cache, heavy vectorized operations may cause the DataFrame to be moved from the ultra-fast cache to the slower memory. Using eval() and query() are potentially more efficient in those cases, but even then the gain in performance and saving in memory is marginal.

The benefit becomes more obvious for big datasets (gigabytes). The intermediate steps create full copies of the underlying DataFrame, so that the data may not even fit into the memory and needs to be swapped on the disk. The computation might not even terminate if the computer runs out of swap space. In those cases, eval() and query() not only help saving memory, but also make some operations possible in the first place.

3.14 Miscellaneous

Pandas allows to read CSV files into a DataFrame. Given the CSV file countries.csv, it can be read as follows:


```
Country, Population, Area
USA, 326625792, 9147593
Russia, 142257520, 16377742
Germany, 80594016, 348672
Switzerland, 8236303, 39997
```

```
>>> countries = pd.read_csv('countries.csv')
>>> countries
   Country  Population    Area
0      USA   326625792  9147593
1    Russia  142257520 16377742
2   Germany   80594016   348672
3 Switzerland   8236303    39997
```

Data can also be read from JSON files, like `countries.json`, which can be read as follows:

```
{
  "country": [
    "USA",
    "Russia",
    "Germany",
    "Switzerland"
  ],
  "population": [
    326625792,
    142257520,
    80594016,
    8236303
  ],
  "area": [
    9147593,
    16377742,
    348672,
    39997
  ]
}

>>> countries = pd.read_json('countries.json')
>>> countries
   country  population    area
0      USA   326625792  9147593
1    Russia  142257520 16377742
2   Germany   80594016   348672
3 Switzerland   8236303    39997
```

4 Matplotlib

Matplotlib is a multiplatform data visualization library built on NumPy arrays. It supports different graphic backends and output styles, and works on virtually any platform. Some projects, including Pandas, offer wrappers around the API of Matplotlib. It is, however, still useful to know how to deal directly with Matplotlib.

Conventionally, Matplotlib is imported as follows:

```
>>> import matplotlib as mpl
>>> import matplotlib.pyplot as plt
```

The plot style can be set on the `plt` object:

```
>>> plt.style.use('classic')
```

Depending on the context, there are different ways of opening the plots for display.

From a script, the method `plt.show` opens all figures plotted so far:

```
import matplotlib as mpl
import matplotlib.pyplot as plt
import numpy as np

x = np.linspace(0, 10, 100)
plt.plot(x, np.sin(x))
plt.plot(x, np.cos(x))

plt.show()
```

The method `plt.show` must only be used once per script or session.

Plots created in a IPython shell can be displayed automatically by calling the `%matplotlib` magic command before calling methods on the `plt` object. The plot will be displayed in a separate window. The method `plt.draw` forces the output to be updated.

```
>>> import matplotlib as mpl
>>> import matplotlib.pyplot as plt
>>> import numpy as np

>>> %matplotlib
Using matplotlib backend: Qt5Agg

>>> x = np.linspace(0, 10, 100)
>>> plt.plot(x, np.sin(x))
```

From within a Jupyter Notebook, there are two options to display plots:

1. `%matplotlib inline`: display plots as static images
2. `%matplotlib notebook`: display interactive plots

The latter option will draw every plot output in the most recent figure, which can be created using the `plt.figure` method:

```
import matplotlib as mpl
import matplotlib.pyplot as plt
import numpy as np

x = np.linspace(0, 10, 100)

%matplotlib notebook

plt.figure()
plt.plot(x, np.sin(x))
plt.plot(x, np.cos(x))
```

A figure object can be saved using its `savefig` method, which requires a file name. Notice that the plot method only draws into the most recent figure object created, if the magic command `%matplotlib` hasn't been used before:

```
import matplotlib as mpl
import matplotlib.pyplot as plt
import numpy as np

fig = plt.figure()
x = np.linspace(0, 10, 100)
plt.plot(x, np.sin(x))
plt.plot(x, np.cos(x))
fig.savefig('sin-x-cos-x.png')
```




Figure 5: Plot of $\sin(x)$ and $\cos(x)$

An image-no longer a plot!-can be loaded using IPython's Image object:

```
>>> from IPython.display import Image, display
>>> img = Image('sin-x-cos-x.png')
>>> display(img)
```

For both saving and loading, the file format is inferred from the file's extension. The formats supported by the graphics backend in use can be retrieved as a dictionary from a figure object:

```
>>> import matplotlib as mpl
>>> import matplotlib.pyplot as plt

>>> fig = plt.figure()
>>> fig.canvas.get_supported_filetypes()
```

```
{'ps': 'Postscript',
 'eps': 'Encapsulated Postscript',
 'pdf': 'Portable Document Format',
 'pgf': 'PGF code for LaTeX',
 'png': 'Portable Network Graphics',
 'raw': 'Raw RGBA bitmap',
 'rgba': 'Raw RGBA bitmap',
 'svg': 'Scalable Vector Graphics',
 'svgz': 'Scalable Vector Graphics'}
```

4.1 Interfaces: MATLAB-style and Object Oriented

Matplotlib started out as a Python alternative for MATLAB. The `plt` object represents the stateful interface known to MATLAB users. Plots created on the `plt` object are drawn to the figure and axes objects that have been created most recently.

In this example, two subplots on a single figure are created:

```
import matplotlib.pyplot as plt
import numpy as np

x = np.linspace(0, 10, 100)
plt.figure() # create a new figure
plt.subplot(2, 1, 1) # (row, column, panel): first panel on a 2*1 field
plt.plot(x, np.sin(x)) # plot to the first subplot
plt.subplot(2, 1, 2) # second panel on the same 2*1 field
plt.plot(x, np.cos(x)) # plot to the second subplot
plt.show()
```

It is possible to plot on other figures/axes than the current active, but only if their references have been retrieved and stored using `plt.gcf` (get current figure) and `plt.gca` (get current axes):

```
import matplotlib.pyplot as plt
import numpy as np

x = np.linspace(0, 10, 100)
plt.figure()
plt.subplot(2, 1, 1)
plt.plot(x, np.sin(x))
first = plt.gca() # store reference to first axes
plt.subplot(2, 1, 2)
plt.plot(x, np.cos(x))
first.plot(x, np.cos(x)) # also draw cosine on first axes
plt.show()
```

“Going back” is not possible if one fails to store the such references, especially in an interactive session. The object-oriented interface of Matplotlib doesn’t rely on a *current state*, but requires the user to always explicitly refer to the figure/axes to be dealt with:

```
import matplotlib.pyplot as plt
import numpy as np

x = np.linspace(0, 10, 100)
fig, ax = plt.subplots(2)
ax[0].plot(x, np.sin(x))
```




Figure 6: MATLAB-style interface: Subplots




Figure 7: MATLAB-style interface: Draw to “inactive” Axes

```
ax[1].plot(x, np.cos(x))
ax[0].plot(x, np.cos(x))
plt.show()
```

The choice between the two interfaces is mostly a matter of preference for simple tasks. More complicated plots, however, do require the object-oriented approach.

4.2 Line Plots

Simple functions of the form $y=f(x)$ can be visualized using line plots. The following examples require this boilerplate code:

```
import matplotlib.pyplot as plt
import numpy as np

%matplotlib inline

plt.style.use('seaborn-whitegrid') # simple style

x = np.linspace(0, 10, 100) # 100 points in range 0..10
```

A figure, implemented by `plt.Figure`, contains all the graphics objects, like text, labels – and the axes. A axes, implemented by `plt.Axes`, is a bounding box with ticks and labels, which contains the plotted lines. Conventionally, the objects are called `fig` and `ax`:

```
fig = plt.figure()
ax = plt.axes()
```

The sine function of the `x` values computed before can be drawn using the `axes` plot method:

```
ax.plot(x, np.sin(x)) # plot x and y=sin(x)
```

The MATLAB-style interface can be used alongside, plotting to the figure/axes used most recently:

```
plt.plot(x, np.cos(x)) # plot x and y=cos(x)
```

The lines get a color assigned automatically from a predefined set. The colors can also be assigned manually using the `color` keyword of the `plot` method. The following options are supported:

- HTML color name: blue, green, fuchsia etc. (common HTML color names)
- RGB/CMYK short code: r, g, b, c, m, y, k
- Grayscale value: floating point number between 0 (black) and 1 (white)
- RGB hex code: #ff00aa, #efefef
- RGB tuple with floating point numbers between 0 and 1: (0.1, 0.75, 0.66)

The line style can be adjusted using the `linestyle` keyword. The following options are supported, both having a short and a long form:

- `-/solid`
- `--/dashed`
- `-. /dashdot`
- `:/dotted`

```

ax.plot(x, np.sin(x-0), color='fuchsia', linestyle='-')
ax.plot(x, np.sin(x-1), color='m', linestyle='--')
ax.plot(x, np.sin(x-2), color='0.25', linestyle='-.')
ax.plot(x, np.sin(x-3), color='#0a123b', linestyle=':')
ax.plot(x, np.sin(x-4), color=(0.1, 0.75, 0.66))

```




Figure 8: Line Colors and Styles

The MATLAB-style interface `plt.plot` accepts a shorthand style indicator as a third non-keyword argument, combining a line style with a RGB/CMYK color code:

```
plt.plot(x, np.sin(x-5), ':y') # dotted yellow line
```


4.3 Limits, Labels, Legends

The axes limits can be set using the `plt.xlim` and `plt.ylim` function (MATLAB-style) or using the `ax.set_xlim` and `ax.set_ylim` method (OO-style) by passing a lower and a upper bound:

```
plt.xlim(-5, 5)
plt.ylim(-1, 1)
```

A plot can be flipped along both axis by passing the arguments in reverse order (using an axes object here):

```
ax.set_xlim(5, -5)
ax.set_ylim(1, -1)
```

The `plt.axis` method allows to set both ranges at once by providing a list of the form `[xmin, xmax, ymin, ymax]`:

```
plt.axis([-5, 5, -1, 1])
```

The ranges can be set automatically to just fit in the plot by using the `'tight'` parameter:

```
plt.axis('tight')
```

The `'equal'` parameter makes sure the plot fits in and that the x and y axis are scaled equally:

```
plt.axis('equal')
```

The axes object supports the same method: `ax.axis`.

Both axis and the plot as a whole can be labeled using the `plt.xlabel`, `plt.ylabel` and `plt.title` function (MATLAB-style) or the `ax.set_xlabel`, `ax.set_ylabel`, `ax.set_title` method (OO-style):

```
plt.xlabel('x')
plt.ylabel('y=sin(x)')
plt.title('A Sine Curve')
```

```
ax.set_xlabel('x')
ax.set_ylabel('y=cos(x)')
ax.set_title('A Cosine Curve')
```

Lines with different styles and colors can be labeled with a legend by calling the `plt.legend` function or the `ax.legend` method, which requires the individual plots (as opposed to its axis) to be labeled with the plot call (keyword `label`):

```
plt.plot(x, np.sin(x), '-g', label='sin(x)')
plt.plot(x, np.cos(x), ':b', label='cos(x)')
plt.legend()
```

```
ax.plot(x, np.sin(x), color='green', linestyle='-', label='sin(x)')
ax.plot(x, np.cos(x), color='blue', linestyle=':', label='cos(x)')
ax.legend()
```

The `ax.set` method is a convenient interface for setting limits (using tuples), labels and a title all at once:

```
ax.set(xlim=(0, 10), ylim=(-1, 1),
       xlabel='x', ylabel='sin(x)',
       title='A Sine Curve')
```

Bringing it all together (in a script):

```
import matplotlib.pyplot as plt
import numpy as np

x = np.linspace(0, 10, 100)

ax = plt.axes()
ax.set(xlim=(0, 10), ylim=(-1, 1),
       xlabel='x', ylabel='y',
       title='Sine and Cosine')

ax.plot(x, np.sin(x), color='green', linestyle='-', label='sin(x)')
ax.plot(x, np.cos(x), color='blue', linestyle=':', label='cos(x)')
ax.legend()

plt.show()
```

4.4 Scatter Plots

Scatter plots represent the data points individually instead of joining them with a line. The `plt.plot` function is capable of producing scatter plots, if the third argument is a character representing an according symbol, such as 'o', '.', ',', 'x', '+', 'v', '^', '<', '>', 's', 'd':

```
import matplotlib.pyplot as plt
import numpy as np

x = np.linspace(0, 10, 50)
plt.plot(x, np.sin(x), 'o', label='sin(x)')
plt.plot(x, np.cos(x), 'x', label='cos(x)')

plt.legend()
plt.show()
```

The dots can be connected when combining the style parameter with a line style:

```
plt.plot(x, np.sin(x), 'o-c', label='sin(x)') # dots & solid cyan line
plt.plot(x, np.cos(x), 'x:m', label='cos(x)') # crosses & dotted magenta line
```

The lines and markers (points) can be further specified using the following arguments of the `plt.plot` function:

- `markersize`
- `markerfacecolor`
- `markeredgecolor`
- `markeredgewidth`
- `linewidth`

```
plt.plot(x, np.sin(x), 'o-c', markersize=10, markerfacecolor='blue',
         markeredgecolor='white', markeredgewidth=2, linewidth=3)
```




Figure 9: Limits, Labels, Legends




Figure 10: Sine and Cosine curve as a scatter plot




Figure 11: Points connected with a line




Figure 12: Marker options

The function `plt.scatter` can set the individual properties of each point by passing a list instead of a single value as the size and color parameters (s and c):

```
n_points = 100
rng = np.random.RandomState(0)
x = rng.randn(n_points)
y = rng.randn(n_points)

colors = rng.randn(n_points)
sizes = 500 * rng.randn(n_points)

plt.scatter(x, y, c=colors, s=sizes, alpha=0.3)

plt.colorbar()
plt.show()
```

This is useful for visualizing multi-dimensional data (four dimensions: x and y value, color and size).

Because `plt.scatter` figures out the rendering for each individual point separately, it can be slower than `plt.plot`, especially when dealing with big data sets. If all the scatter points are to be drawn alike, `plt.plot` should be preferred to `plt.scatter`.

The OO-style interface (`ax.plot`, `ax.scatter`) works with the same parameters.

4.5 Visualizing Errors

In many applications, reporting the range of possible error is just as important as reporting the value itself.

For discrete values, Matplotlib can plot error bars using the `plt.errorbar` function. The error range, either on the x- or y-axis, can be set using the parameter `xerr` or `yerr`, respectively. The `fmt` parameter accepts a format specifier consisting of style and color code:

```
rng = np.random.RandomState(0)
points = 20
dy = 0.5

x = np.linspace(0, 10, points)
err = dy * rng.randn(points)
y = np.sin(x) + err

plt.errorbar(x, y, yerr=dy, fmt='.')
plt.show()
```

The error bar can be further fine-tuned by specifying the `ecolor` (color of the bar), the `elinewidth` (the width of the error bar) and the `capsize` (the size of the ticks orthogonal to the error bar) parameters:

```
plt.errorbar(x, y, yerr=dy, fmt='.', ecolor='#cccccc', elinewidth=2, capsize=5)
```

The error of continuous quantities can be indicated by filling a area around the graph displaying the values. This can be achieved by using a combination of the `plt.plot` (indicating the values) and the `plt.fill_between` (indicating the area of error) function.




Figure 13: Scatter plot with individual marker sizes and colors

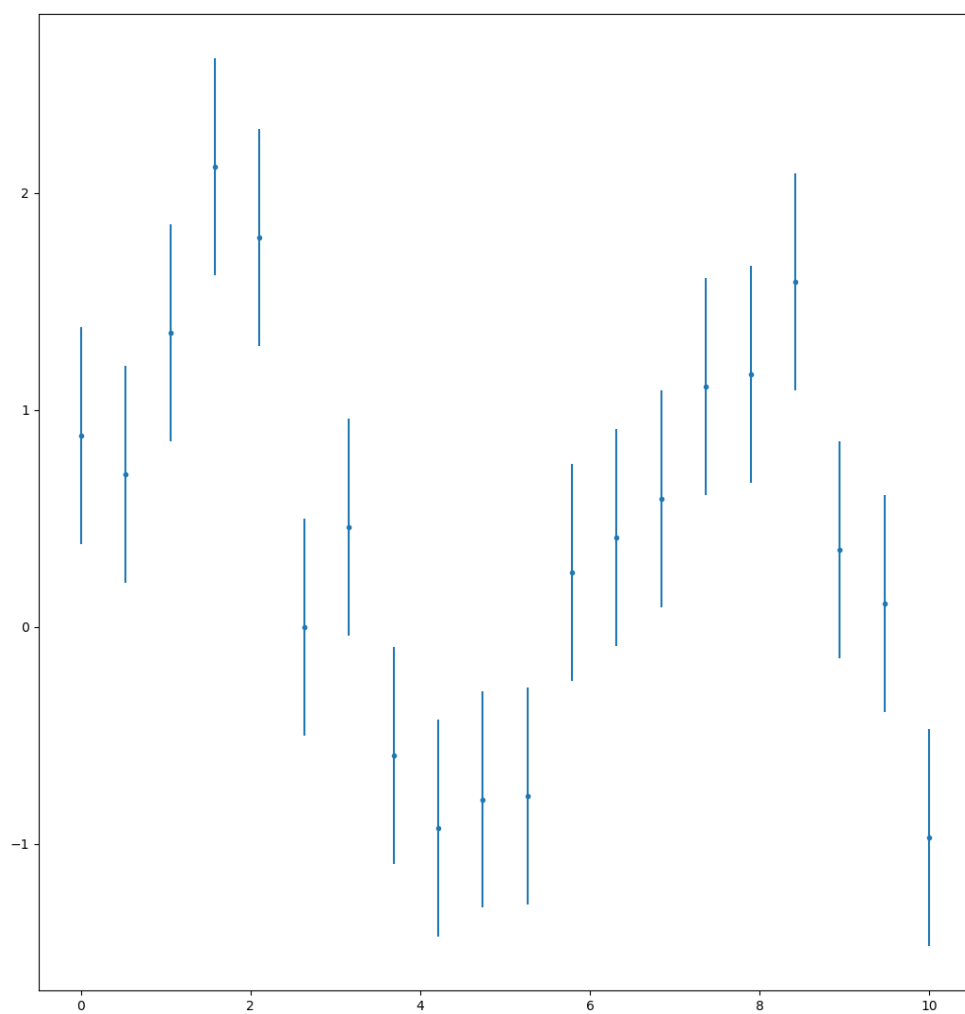


Figure 14: Vertical Error Bars




Figure 15: Customized Error Bars

```

rng = np.random.RandomState(0)
points = 100
dy = 0.25

x = np.linspace(0, 10, points)
err = dy * rng.randn(points)
y = np.sin(x) + err

plt.plot(x, y)
plt.fill_between(x, y-dy, y+dy, color='#cccccc', alpha=.5)
plt.show()

```

The second argument ($y-dy$) is the lower, the third argument ($y+dy$) the upper bound of the error area.

The methods `errorbar` and `fill_between` are also available in the axes' OO-style interface:

```

fig, ax = plt.subplots(2)
ax[0].errorbar(x, y, yerr=err, fmt='.')
ax[1].plot(x, y)
ax[1].fill_between(x, y-dy, y+dy, color='#cccccc', alpha=.5)
plt.show()

```

4.6 Density and Contour Plots

Three-dimensional data can be displayed in two dimensions using contours or color-coded regions. A function $z=f(x, y)$ can be visualized by using x and y as the positions on the grid, and z for the contour level:

```

def f(x, y):
    return np.sin(x) + np.cos(x * y) * np.cos(x)

```

The z values are broadcasted into a two-dimensional grid. For the x and y values, broadcasting can be done using the `np.meshgrid` function:

```

x = np.linspace(0, 5, 50)
y = np.linspace(0, 5, 40)
X, Y = np.meshgrid(x, y)
Z = f(X, Y)

```

The contour plot can be created using the `plt.contour` function:

```

plt.contour(X, Y, Z, colors='black')
plt.show()

```

Negative z values are represented by dashed, positive z values by solid lines. A color code with a number of intervals can be used in conjunction with a colormap instead:

```

plt.contour(X, Y, Z, 20, cmap='RdGy')

```

`RdGy` is a red-green colormap, with red indicating negative and grey positive values. More colormaps are available under `plt.cm`.

Instead of a contour plot with its distracting gaps, a *filled* contour plot can be created using the `plt.contourf` function:

```

plt.contourf(X, Y, Z, 20, cmap='RdGy')
plt.colorbar()

```




Figure 16: Error Area for Continuous Values




Figure 17: Error Bar and Area Combined



Figure 18: Contour Plot of a $z=f(x,y)$ Function




Figure 19: Contour Plot Using a Color Map




Figure 20: Filled Contour Plot

The colorbar helps to identify peaks and valleys. The color steps are discrete (20 contours) rather than continuous. The number of contours could be increased, which would be rather inefficient. The `plt.imshow` function is a faster option for that purpose:

```
plt.imshow(Z, extent=[0, 5, 0, 5], origin='lower', cmap='RdGy')
plt.colorbar()
plt.axis(aspect='image')
```




Figure 21: Contour Plot as an Image

- Instead of a grid, the value range on the x and y axis is defined as the extent of the form `[xmin, xmax, ymin, ymax]`.
- The image is drawn from the lower left (like a function), not from the upper left (like an image) by setting the origin argument lower.

- To match the x and y units, the aspect argument is set to image to prevent automatic aspect ratio adjustment.

Image and contour plots can also be combined, and the contours can also be labeled with their value:

```
contours = plt.contour(X, Y, Z, 3, colors='black')
plt.clabel(contours, inline=True, fontsize=8)
plt.imshow(Z, extent=[0, 5, 0, 5], origin='lower', cmap='RdGy', alpha=0.5)
plt.colorbar()
```




Figure 22: Partially Transparent Background Image, Over-Plotted with Contours and Labels

4.7 Histograms, Binnings and Density

A histogram is a good first step in understanding a dataset. The function `plt.hist` creates a histogram of a one-dimensional array of numeric values:

```
import numpy as np
import matplotlib.pyplot as plt

data = np.random.randn(1000)
plt.hist(data, color='lightblue', edgecolor='black')

plt.show()
```




Figure 23: Histogram of 1000 Normal Distributed Values

Multiple datasets can be easily compared by plotting them as transparent histograms:

```
x1 = np.random.normal(-1, 0.75, 1000)
x2 = np.random.normal(0, 1, 1000)
x3 = np.random.normal(1, 1.25, 1000)

kwargs = dict(alpha=0.3, bins=50, edgecolor='black', histtype='stepfilled')
plt.hist(x1, **kwargs, color='#ff0000')
plt.hist(x2, **kwargs, color='#00ff00')
plt.hist(x3, **kwargs, color='#0000ff')

plt.show()
```




Figure 24: Histograms of 3 Normal Distributed Datasets

The commonly used named arguments (such as `histtype='stepfilled'`, which draws a contour around the overall histogram area, as opposed to each separate bin) are passed as *variable length keyworded arguments* using the `**kwargs` syntax. The dataset-specific color argument is defined and passed individually for each dataset, in order to distinguish them on the plot.

The `np.histogram` function allows to *calculate* a histogram without drawing it, returning a tuple of arrays for the quantities (y axis) and bin locations (x axis):

```
>>> data = np.random.randn(1000)
>>> np.histogram(data, bins=10)
(array([ 9, 28, 75, 171, 200, 228, 160, 83, 39, 7]),
 array([-2.91074477, -2.33086085, -1.75097692, -1.171093, -0.59120908,
        -0.01132516, 0.56855876, 1.14844269, 1.72832661, 2.30821053,
        2.88809445]))
```